



Microsoft Cloud Workshop

Cloud-native applications - Infrastructure edition

Hands-on lab step-by-step

August 2020

Information in this document, including URL and other Internet Web site references, is subject to change without notice. Unless otherwise noted, the example companies, organizations, products, domain names, e-mail addresses, logos, people, places, and events depicted herein are fictitious, and no association with any real company, organization, product, domain name, e-mail address, logo, person, place or event is intended or should be inferred. Complying with all applicable copyright laws is the responsibility of the user. Without limiting the rights under copyright, no part of this document may be reproduced, stored in or introduced into a retrieval system, or transmitted in any form or by any means (electronic, mechanical, photocopying, recording, or otherwise), or for any purpose, without the express written permission of Microsoft Corporation.

Microsoft may have patents, patent applications, trademarks, copyrights, or other intellectual property rights covering subject matter in this document. Except as expressly provided in any written license agreement from Microsoft, the furnishing of this document does not give you any license to these patents, trademarks, copyrights, or other intellectual property.

The names of manufacturers, products, or URLs are provided for informational purposes only and Microsoft makes no representations and warranties, either expressed, implied, or statutory, regarding these manufacturers or the use of the products with any Microsoft technologies. The inclusion of a manufacturer or product does not imply endorsement of Microsoft of the manufacturer or product. Links may be provided to third party sites. Such sites are not under the control of Microsoft and Microsoft is not responsible for the contents of any linked site or any link contained in a linked site, or any changes or updates to such sites. Microsoft is not responsible for webcasting or any other form of transmission received from any linked site. Microsoft is providing these links to you only as a convenience, and the inclusion of any link does not imply endorsement of Microsoft of the site or the products contained therein.

© 2020 Microsoft Corporation. All rights reserved.

Microsoft and the trademarks listed at <https://www.microsoft.com/en-us/legal/intellectualproperty/Trademarks/Usage/General.aspx> are trademarks of the Microsoft group of companies. All other trademarks are property of their respective owners.

Contents

- [Cloud-native applications - Infrastructure edition hands-on lab step-by-step](#)
 - [Abstract and learning objectives](#)
 - [Overview](#)
 - [Solution architecture](#)
 - [Requirements](#)
 - [Exercise 1: Create and run a Docker application](#)
 - [Task 1: Test the application](#)
 - [Task 2: Browsing to the web application](#)
 - [Task 3: Create Docker images](#)
 - [Task 4: Run a containerized application](#)
 - [Task 5: Setup environment variables](#)
 - [Task 6: Run several containers with Docker compose](#)
 - [Task 7: Push images to Azure Container Registry](#)
 - [Task 8: Setup CI Pipeline to Push Images](#)

- Exercise 2: Deploy the solution to Azure Kubernetes Service
- Task 1: Tunnel into the Azure Kubernetes Service cluster
- Task 2: Deploy a service using the Kubernetes management dashboard
- Task 3: Deploy a service using kubectl
- Task 4: Deploy a service using a Helm chart
- Task 5: Initialize database with a Kubernetes Job
- Task 6: Test the application in a browser
- Task 7: Configure Continuous Delivery to the Kubernetes Cluster
- Task 8: Review Azure Monitor for Containers
- Exercise 3: Scale the application and test HA
- Task 1: Increase service instances from the Kubernetes dashboard
- Task 2: Increase service instances beyond available resources
- Task 3: Restart containers and test HA
- Exercise 4: Working with services and routing application traffic
- Task 1: Scale a service without port constraints
- Task 2: Update an external service to support dynamic discovery with a load balancer
- Task 3: Adjust CPU constraints to improve scale
- Task 4: Perform a rolling update
- Task 5: Configure Kubernetes Ingress
- After the hands-on lab

Cloud-native applications - Infrastructure edition hands-on lab step-by-step

Abstract and learning objectives

This hands-on lab is designed to guide you through the process of building and deploying Docker images to the Kubernetes platform hosted on Azure Kubernetes Services (AKS), in addition to learning how to work with dynamic service discovery, service scale-out, and high-availability.

At the end of this lab you will be better able to build and deploy containerized applications to Azure Kubernetes Service and perform common DevOps procedures.

Overview

Fabrikam Medical Conferences (FabMedical) provides conference website services tailored to the medical community. They are refactoring their application code, based on node.js, so that it can run as a Docker application, and want to implement a POC that will help them get familiar with the development process, lifecycle of deployment, and critical aspects of the hosting environment. They will be deploying their applications to Azure Kubernetes Service and want to learn how to deploy containers in a dynamically load-balanced manner, discover containers, and scale them on demand.

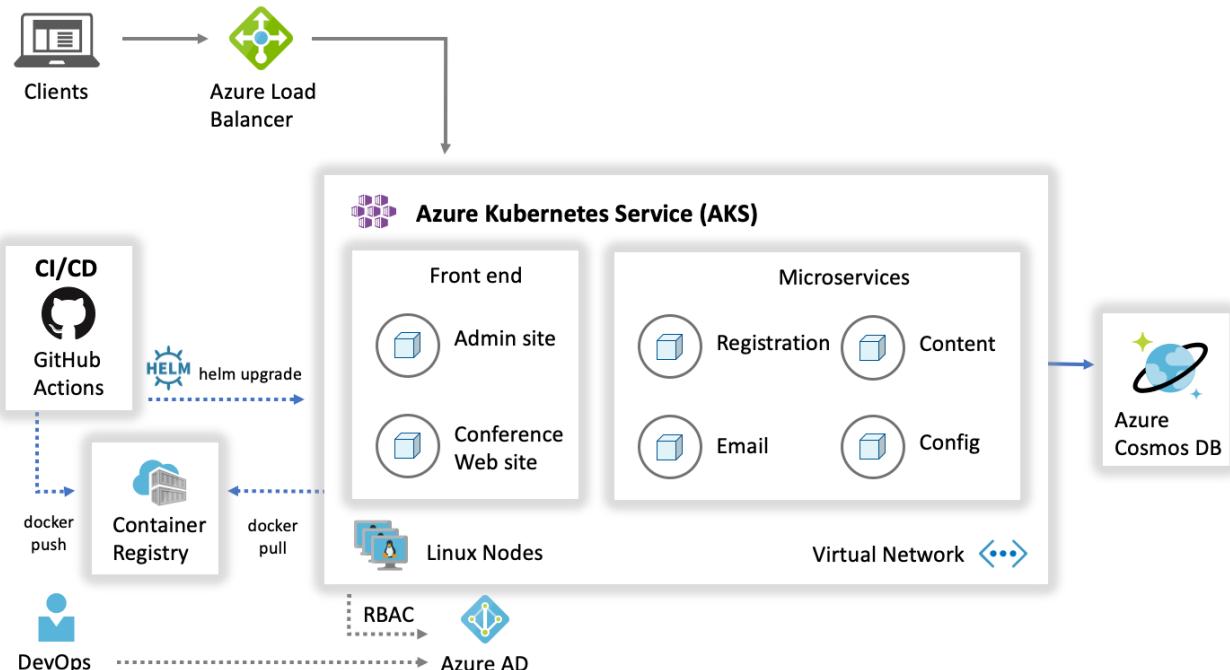
In this hands-on lab, you will assist with completing this POC with a subset of the application codebase. You will create a build agent based on Linux, and an Azure Kubernetes Service cluster for running deployed applications. You will be helping them to complete the Docker setup for their application, test locally, push to an image repository, deploy to the cluster, and test load-balancing and scale.

Important: Most Azure resources require unique names. Throughout these steps, you will see the word "SUFFIX" as part of resource names. You should replace this with a unique handle (like your Microsoft Account email prefix) to ensure unique names for resources.

Solution architecture

Below is a diagram of the solution architecture you will build in this lab. Please study this carefully, so you understand the whole of the solution as you are working on the various components.

The solution will use Azure Kubernetes Service (AKS), which means that the container cluster topology is provisioned according to the number of requested nodes. The proposed containers deployed to the cluster are illustrated below with Cosmos DB as a managed service:



Each tenant will have the following containers:

- **Conference Web site:** The SPA application that will use configuration settings to handle custom styles for the tenant.
- **Admin Web site:** The SPA application that conference owners use to manage conference configuration details, manage attendee registrations, manage campaigns, and communicate with attendees.
- **Registration service:** The API that handles all registration activities creating new conference registrations with the appropriate package selections and associated cost.
- **Email service:** The API that handles email notifications to conference attendees during registration, or when the conference owners choose to engage the attendees through their admin site.
- **Config service:** The API that handles conference configuration settings such as dates, locations, pricing tables, early-bird specials, countdowns, and related.
- **Content service:** The API that handles content for the conference such as speakers, sessions, workshops, and sponsors.

Requirements

1. Microsoft Azure subscription must be pay-as-you-go or MSDN.
 - Trial subscriptions will *not* work.
 - To complete this lab, ensure your account has the following roles:
 - The [Owner](#) built-in role for the subscription you will use.

- Is a [Member](#) user in the Azure AD tenant you will use. (Guest users will not have the necessary permissions).

Note If you do not meet these requirements, you may have to ask another member user with subscription owner rights to login to the portal and execute the create service principal step ahead of time.

- You must have enough cores available in your subscription to create the build agent and Azure Kubernetes Service cluster in Before the Hands-on Lab. You will need eight cores if following the exact instructions in the lab, or more if you choose additional cluster nodes or larger VM sizes. If you execute the steps required before the lab, you will be able to see if you need to request more cores in your sub.

2. Local machine or a virtual machine configured with:

- A browser, preferably Chrome for consistency with the lab implementation tests.

3. You will install other tools throughout the exercises.

Very important: You should be typing all the commands as they appear in the guide. Do not try to copy and paste to your command windows or other documents when instructed to enter the information shown in this document, except where explicitly stated in this document. There can be issues with Copy and Paste that result in errors, execution of instructions, or creation of file content.

Exercise 1: Create and run a Docker application

Duration: 40 minutes

In this exercise, you will take the starter files and run the node.js application as a Docker application. You will create a Dockerfile, build Docker images, and run containers to execute the application.

Task 1: Test the application

The purpose of this task is to make sure you can run the application successfully before applying changes to run it as a Docker application.

1. From Azure Cloud Shell, connect to your build agent if you are not already connected. (If you need to reconnect, please review the instructions in the "Before the HOL" document.)
2. Type the following command to create a Docker network named `fabmedical`:

```
docker network create fabmedical
```

3. Run an instance of mongodb to use for local testing.

```
docker container run --name mongo --net fabmedical -p 27017:27017 -d mongo
```

Note: With the existing source code written for MongoDB, it can be pointed towards the Azure Cosmos DB MongoDB API endpoint. The Azure Cosmos DB Emulator could be used for local development on Windows, however, the Cosmos DB emulator does not support Linux. As a result, when using Linux for development, MongoDB is still needed for local development environments; with Azure Cosmos DB used for data storage in the cloud. This allows existing source code written for MongoDB storage to be easily migrated to using Azure Cosmos DB backend.

4. Confirm that the mongo container is running and ready.

```
docker container list  
docker container logs mongo
```

5. Connect to the mongo instance using the mongo shell and test some basic commands:

mongo

```
show dbs  
quit()
```

```
Bash ~ | ⌂ ? ⌂ {} ⌂
admin@fabmedical:~$ mongo
MongoDB shell version: 2.6.10
connecting to: test
Welcome to the MongoDB shell.
For interactive help, type "help".
For more comprehensive documentation, see
    http://docs.mongodb.org/
Questions? Try the support group
    http://groups.google.com/group/mongodb-user
Server has startup warnings:
2019-10-04T12:45:09.692+0000 I STORAGE [initandlisten]
2019-10-04T12:45:09.692+0000 I STORAGE [initandlisten] ** WARNING: Using the XFS filesystem is strongly recommended with the WiredTiger storage engine
2019-10-04T12:45:09.692+0000 I STORAGE [initandlisten] ** See http://dochub.mongodb.org/core/prodnotes-filesystem
2019-10-04T12:45:10.630+0000 I CONTROL [initandlisten]
2019-10-04T12:45:10.630+0000 I CONTROL [initandlisten] ** WARNING: Access control is not enabled for the database.
2019-10-04T12:45:10.630+0000 I CONTROL [initandlisten] ** Read and write access to data and configuration is unrestricted.
2019-10-04T12:45:10.630+0000 I CONTROL [initandlisten]
2019-10-04T12:45:10.630+0000 I CONTROL [initandlisten]
2019-10-04T12:45:10.630+0000 I CONTROL [initandlisten] ** WARNING: /sys/kernel/mm/transparent_hugepage/enabled is 'always'.
2019-10-04T12:45:10.630+0000 I CONTROL [initandlisten] ** We suggest setting it to 'never'
2019-10-04T12:45:10.630+0000 I CONTROL [initandlisten]
> show dbs
admin 0.000GB
config 0.000GB
local 0.000GB
> quit()
admin@fabmedical:~$
```

6. To initialize the local database with test content, first navigate to the content-init directory and run npm install.

```
cd content-init  
npm install
```

Note: In some cases, the `root` user will be assigned ownership of your user's `.config` folder. If this happens, run the following command to return ownership to `adminfabmedical` and then try `npm install` again:

```
sudo chown -R $USER:$(id -gn $USER) /home/adminfabmedical/.config
```

7. Initialize the database.

nodejs server.js

```
Bash    v | ⌂ ? ⌂ ⌂ ⌂ {} ⌂
adminfabmedical@fabmedical-:~$ cd content-init/
adminfabmedical@fabmedical-:~/content-init$ npm install
npm WARN content-init@1.0.0 No description
npm WARN content-init@1.0.0 No repository field.

added 28 packages from 17 contributors and audited 35 packages in 1.61s
found 0 vulnerabilities

adminfabmedical@fabmedical-:~/content-init$ nodejs server.js
Clean Sessions table
(node:9955) DeprecationWarning: collection.remove is deprecated. Use deleteOne, deleteMany, or bulkWrite instead
Connected to MongoDB
All Sessions deleted
Load sessions from JSON file
Session saved successfully
Session saved successfully
Session saved successfully
Session saved successfully
Clean Speakers table
All Speakers deleted
Load Speakers from JSON file
Speaker saved successfully
Speaker saved successfully
Speaker saved successfully
Speaker saved successfully
adminfabmedical@fabmedical-:~/content-init$
```

8. Confirm that the database now contains test data

mongo

```
show dbs
use contentdb
show collections
db.speakers.find()
db.sessions.find()
quit()
```

This should produce output similar to the following:



9. Now navigate to the `content-api` directory and run `npm install`

```
cd ..../content-api  
npm install
```

Note: In some cases, the root user will be assigned ownership of your user's .config folder. If this happens, run the following command to return ownership to adminfabmedical and then try npm install again:

```
sudo chown -R $USER:$(id -gn $USER) /home/adminfabmedical/.config
```

10. Start the API as a background process.

```
nodejs ./server.js &
```

```
Bash └─ ? { } 
adminfabmedical@fabmedical- [~]:~/content-init$ cd ../content-api/
adminfabmedical@fabmedical- [~]:~/content-api$ npm install
added 75 packages from 49 contributors and audited 190 packages in 2.447s
found 0 vulnerabilities

adminfabmedical@fabmedical- [~]:~/content-api$ nodejs ./server.js &
[1] 11912
adminfabmedical@fabmedical- [~]:~/content-api$ Listening on port 3001
Connected to MongoDB

adminfabmedical@fabmedical- [~]:~/content-api$
```

11. Press **ENTER** again to get to a command prompt for the next step.

12. Test the API using curl. You will request the speaker's content, and this will return a JSON result.

```
curl http://localhost:3001/speakers
```

13. Navigate to the web application directory, run `npm install` and `ng build`.

```
cd ..../content-web  
npm install  
ng build  
node ./app.js &
```

Azure Cloud Shell

```
Bash    ▾ | ⌂ ? ⓘ ⌂ ⌂ {} ⌂

es/karma/node_modules/fsevents):
npm [WARN] notsup SKIPPING OPTIONAL DEPENDENCY: Unsupported platform for fsevents@1.2.9: wanted {"os":"darwin","arch":"any"} (current: {"os":"linux","arch":"x64"})
npm [WARN] optional SKIPPING OPTIONAL DEPENDENCY: fsevents@1.2.9 (node_modules/@angular/compiler-cli/node_modules/fsevents):
npm [WARN] notsup SKIPPING OPTIONAL DEPENDENCY: Unsupported platform for fsevents@1.2.9: wanted {"os":"darwin","arch":"any"} (current: {"os":"linux","arch":"x64"})
npm [WARN] optional SKIPPING OPTIONAL DEPENDENCY: fsevents@2.0.7 (node_modules/fsevents):
npm [WARN] notsup SKIPPING OPTIONAL DEPENDENCY: Unsupported platform for fsevents@2.0.7: wanted {"os":"darwin","arch":"any"} (current: {"os":"linux","arch":"x64"})

audited 1459 packages in 8.957s
found 780 vulnerabilities (757 low, 5 moderate, 18 high)
  run `npm audit fix` to fix them, or `npm audit` for details
adminfabmedical@fabmedical-cp4:~/Fabmedical/content-web$ ng build
Your global Angular CLI version (10.0.7) is greater than your local
version (8.3.4). The local Angular CLI version is used.

To disable this warning use "ng config -g cli.warnings.versionMismatch fal
se".

chunk {main} main.js, main.js.map (main) 54.3 kB [initial] [rendered]
chunk {polyfills} polyfills.js, polyfills.js.map (polyfills) 282 kB [initi
al] [rendered]
chunk {polyfills-es5} polyfills-es5.js, polyfills-es5.js.map (polyfills-es
5) 602 kB [initial] [rendered]
chunk {runtime} runtime.js, runtime.js.map (runtime) 6.15 kB [entry] [rend
ered]
chunk {styles} styles.js, styles.js.map (styles) 9.72 kB [initial] [render
ed]
chunk {vendor} vendor.js, vendor.js.map (vendor) 3.87 MB [initial] [render
ed]
Date: 2020-08-24T02:16:08.748Z - Hash: ba53090244b74f736cb1 - Time: 11137m
s
adminfabmedical@fabmedical-cp4:~/Fabmedical/content-web$
```

Note: In some cases, the `root` user will be assigned ownership of your user's `.config` folder. If this happens, run the following command to return ownership to `adminfabmedical` and then try `npm install` again:

```
sudo chown -R $USER:$(id -gn $USER) /home/adminfabmedical/.config
```

14. Press `ENTER` again to get a command prompt for the next step.

15. Test the web application using `curl`. You will see HTML output returned without errors.

```
curl http://localhost:3000
```

16. Leave the application running for the next task.

17. If you received a JSON response to the /speakers content request and an HTML response from the web application, your environment is working as expected.

Task 2: Browsing to the web application

In this task, you will browse to the web application for testing.

1. From the Azure portal select the resource group you created named `fabmedical-SUFFIX`.

2. Select the build agent VM named `fabmedical-SUFFIX` from your list of available resources.

NAME	TYPE	LOCATION	...
fabmedical-soll	Virtual machine	East US 2	...
LinuxAsm	Microsoft.Compute/virtualmachines	East US 2	...
fabmedical-soll261	Network interface	East US 2	...
fabmedicalsolldiag273	Storage account	East US 2	...
fabmedicalsolldisks565	Storage account	East US 2	...
fabmedical-soll-ip	Public IP address	East US 2	...
fabmedical-soll-nsg	Network security group	East US 2	...
fabmedical-soll-vnet	Virtual network	East US 2	...

3. From the **Virtual Machine** blade overview, find the **IP address** of the VM.

The screenshot shows the Azure Virtual Machine blade for the 'fabmedical-soll' VM. The left sidebar has 'Overview' selected. The main pane displays basic information: Resource group (fabmedical-), Status (Running), Location (West Europe), Subscription (Standard D2s v3 (2 vcpus, 8 GB memory)), and Subscription ID. On the right, detailed settings are shown, including Computer name (fabmedical-), Operating system (Linux), Size (Standard D2s v3), Public IP address (52.174.141.11, highlighted with a red box), Virtual network/subnet (fabmedical-soll-vnet/default), and DNS name (Configure).

4. Test the web application from a browser. Navigate to the web application using your build agent IP address at port 3000.

```
http://[BUILDAGENTIP]:3000
EXAMPLE: http://13.68.113.176:3000
```

5. Select the Speakers and Sessions links in the header. You will see the pages display the HTML version of the JSON content you curled previously.

6. Once you have verified the application is accessible through a browser, go to your cloud shell window and stop the running node processes.

```
killall nodejs  
killall node
```

Task 3: Create Docker images

In this task, you will create Docker images for the application --- one for the API application and another for the web application. Each image will be created via Docker commands that rely on a Dockerfile.

- From cloud shell connected to the build agent VM, type the following command to view any Docker images on the VM. The list will only contain the mongodb image downloaded earlier.

```
docker image ls
```

- From the `content-api` folder containing the API application files and the new `Dockerfile` you created, type the following command to create a Docker image for the API application. This command does the following:

- Executes the Docker build command to produce the image
- Tags the resulting image with the name `content-api` (`-t`)
- The final dot (`.`) indicates to use the Dockerfile in this current directory context. By default, this file is expected to have the name `Dockerfile` (case sensitive).

```
docker image build -t content-api .
```

- From the `content-init` folder containing the API application files and the new `Dockerfile` you created, type the following command to create a Docker image for the Init application. This command does the following:

- Executes the Docker build command to produce the image
- Tags the resulting image with the name `content-init` (`-t`)
- The final dot (`.`) indicates to use the Dockerfile in this current directory context. By default, this file is expected to have the name `Dockerfile` (case sensitive).

```
docker image build -t content-init .
```

- Once the image is successfully built, run the Docker images listing command again. You will see several new images: the node images and your container image.

```
docker image ls
```

Notice the untagged image. This is the build stage which contains all the intermediate files not needed in your final image.

```
adminfabmedical@fabmedical-:~/Fabmedical/content-api$ docker image ls
REPOSITORY          TAG      IMAGE ID      CREATED       SIZE
content-api         latest   798d89b8f0aa  About a minute ago  129MB
<none>              <none>   8a86289adce3  About a minute ago  674MB
mongo               latest   409c3f937574  4 days ago    493MB
node                alpine   0f2c18cef5d3  11 days ago   117MB
node                argon    ef4b194d8fcf  2 years ago   653MB
adminfabmedical@fabmedical-:~/Fabmedical/content-api$
```

4. Navigate to the `content-web` folder again and list the files. Note that this folder has a Dockerfile.

```
cd ../content-web
11
```

5. View the Dockerfile contents -- which are similar to the file in the API folder. Type the following command:

```
cat Dockerfile
```

Notice that the `content-web` Dockerfile build stage includes additional tools for a front-end Angular application in addition to installing npm packages.

6. Type the following command to create a Docker image for the web application.

```
docker image build -t content-web .
```

8. Navigate to the `content-init` folder again and list the files. Note that this folder already has a Dockerfile.

```
cd ../content-init
11
```

9. View the Dockerfile contents -- which are similar to the file you created previously in the API folder. Type the following command:

```
cat Dockerfile
```

10. Type the following command to create a Docker image for the init application.

```
docker image build -t content-init .
```

11. When complete, you will see seven images now exist when you run the Docker images command.

```
docker image ls
```

```
admin@fabmedical-OptiPlex-5070:~/Fabmedical/content-web$ docker image ls
REPOSITORY      TAG        IMAGE ID      CREATED       SIZE
content-web     latest     6e333b5cd0ae  About a minute ago  354MB
<none>          <none>    440be66d5b98  2 minutes ago   1.31GB
content-api     latest     798d89b8f0aa  10 minutes ago  129MB
<none>          <none>    8a86289adce3  10 minutes ago  674MB
mongo           latest     409c3f937574  4 days ago    493MB
node            alpine    0f2c18cef5d3  11 days ago   117MB
node            dubnium   5ebbf4bb3837  2 weeks ago   911MB
node            dubnium-alpine 8e473595b853  4 weeks ago   83.5MB
node            argon     ef4b194d8fcf  2 years ago   653MB
admin@fabmedical-OptiPlex-5070:~/Fabmedical/content-web$
```

Task 4: Run a containerized application

The web application container will be calling endpoints exposed by the API application container and the API application container will be communicating with mongodb. In this exercise, you will launch the images you created as containers on the same bridge network you created when starting mongodb.

1. Create and start the API application container with the following command. The command does the following:

- Names the container `api` for later reference with Docker commands.
- Instructs the Docker engine to use the `fabmedical` network.
- Instructs the Docker engine to use port `3001` and map that to the internal container port `3001`.
- Creates a container from the specified image, by its tag, such as `content-api`.

```
docker container run --name api --net fabmedical -p 3001:3001 content-api
```

2. The `docker container run` command has failed because it is configured to connect to mongodb using a localhost URL. However, now that `content-api` is isolated in a separate container, it cannot access mongodb via localhost even when running on the same docker host. Instead, the API must use the bridge network to connect to mongodb.

```
> content-api@0.0.0 start /usr/src/app
> node ./server.js

Listening on port 3001
Could not connect to MongoDB!
MongoTimeoutError: Server selection timed out after 30000 ms
npm ERR! code ELIFECYCLE
npm ERR! errno 255
npm ERR! content-api@0.0.0 start: `node ./server.js`
npm ERR! Exit status 255
npm ERR!
npm ERR! Failed at the content-api@0.0.0 start script.
npm ERR! This is probably not a problem with npm. There is likely additional logging output above.

npm ERR! A complete log of this run can be found in:
npm ERR!     /root/.npm/_logs/2019-12-04T22_39_38_815Z-debug.log
```

3. The `content-api` application allows an environment variable to configure the mongodb connection string. Remove the existing container, and then instruct the docker engine to set the environment variable by adding the `-e` switch to the `docker container run` command. Also, use the `-d` switch to run the api as a daemon.

```
docker container rm api
docker container run --name api --net fabmedical -p 3001:3001 -e MONGODB_CONNECTION=mongodb://mongo:27017/contentdb
```

4. Enter the command to show running containers. You will observe that the `api` container is in the list. Use the docker logs command to see that the API application has connected to mongodb.

```
docker container ls
docker container logs api
```

```
adminfabmedical@fabmedical-cp4:~/Fabmedical/content-web$ docker container ls
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS               NAMES
122437f5a5c4        content-api       "docker-entrypoint.s..."   27 seconds ago    Up 26 seconds      0.0.0.0:3001->3001/tcp   api
ef8d527320a1        mongo              "docker-entrypoint.s..."   58 minutes ago   Up 58 minutes     0.0.0.0:27017->27017/tcp   mongo
adminfabmedical@fabmedical-cp4:~/Fabmedical/content-web$ docker container logs api
> content-api@0.0.0 start /usr/src/app
> node ./server.js
Listening on port 3001
Connected to MongoDB
adminfabmedical@fabmedical-cp4:~/Fabmedical/content-web$
```

5. Test the API by curling the URL. You will see JSON output as you did when testing previously.

```
curl http://localhost:3001/speakers
```

6. Create and start the web application container with a similar `docker container run` command -- instruct the docker engine to use any port with the `-P` command.

```
docker container run --name web --net fabmedical -P -d content-web
```

7. Enter the command to show running containers again, and you will observe that both the API and web containers are in the list. The web container shows a dynamically assigned port mapping to its internal container port 3000.

```
docker container ls
```

```
adminfabmedical@fabmedical-cp4:~/Fabmedical/content-web$ docker container ls
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS               NAMES
aa76f2d69b38        content-web       "docker-entrypoint.s..."   18 seconds ago    Up 17 seconds      0.0.0.0:32768->3000/tcp   web
122437f5a5c4        content-api       "docker-entrypoint.s..."   2 minutes ago    Up 2 minutes      0.0.0.0:3001->3001/tcp   api
ef8d527320a1        mongo              "docker-entrypoint.s..."   About an hour ago Up About an hour   0.0.0.0:27017->27017/tcp   mongo
adminfabmedical@fabmedical-cp4:~/Fabmedical/content-web$
```

8. Test the web application by fetching the URL with curl. For the port, use the dynamically assigned port, which you can find in the output from the previous command. You will see HTML output, as you did when testing previously.

```
curl http://localhost:[PORT]/speakers.html
```

Task 5: Setup environment variables

In this task, you will configure the `web` container to communicate with the API container using an environment variable, similar to the way the mongodb connection string is provided to the `api`.

1. From cloud shell connected to the build agent VM, stop and remove the web container using the following commands.

```
docker container stop web
docker container rm web
```

2. Validate that the web container is no longer running or present by using the `-a` flag as shown in this command. You will see that the **web** container is no longer listed.

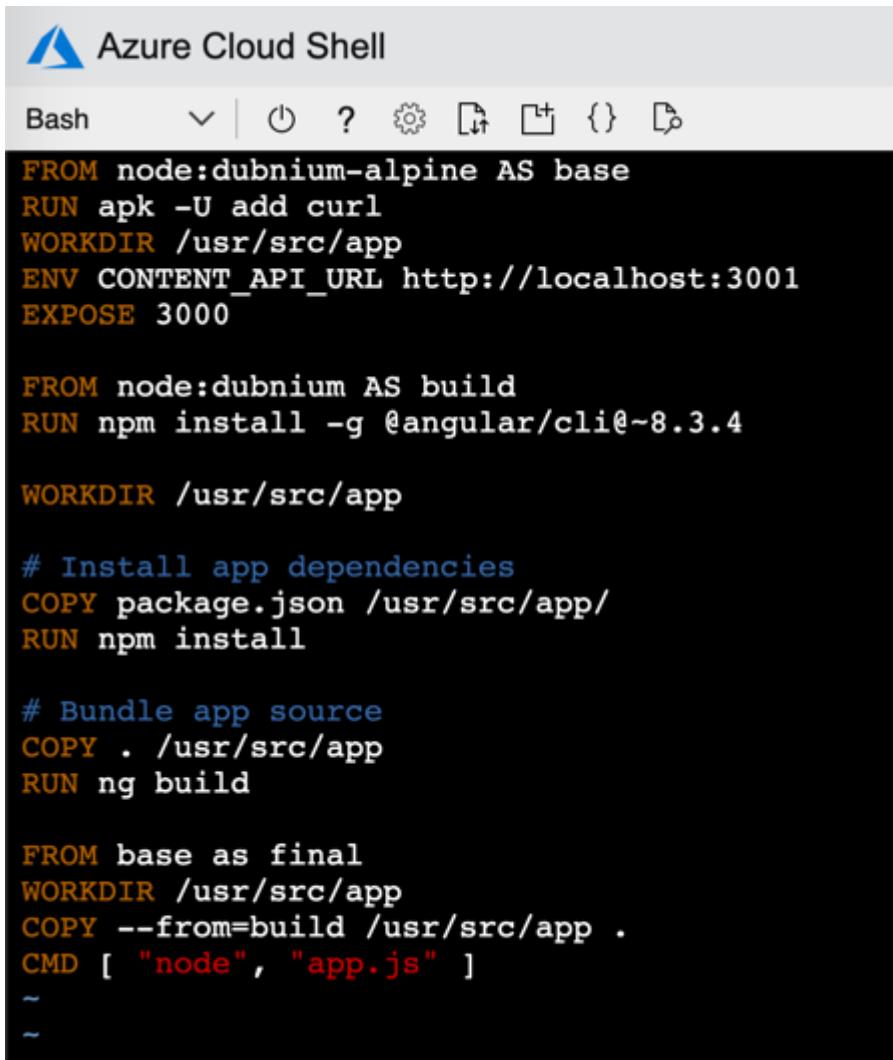
```
docker container ls -a
```

3. Open the Dockerfile for editing using Vim and press the `i` key to go into edit mode.

```
vi Dockerfile  
<i>
```

4. Locate the `EXPOSE` line shown below and add a line above it that sets the default value for the environment variable, as shown in the screenshot.

```
ENV CONTENT_API_URL http://localhost:3001
```



The screenshot shows the Azure Cloud Shell interface with the title "Azure Cloud Shell". Below the title is a toolbar with icons for Bash, dropdown, power, help, settings, copy, paste, and brace matching. The main area is a terminal window displaying a Dockerfile. The Dockerfile contains the following code:

```
FROM node:dubnium-alpine AS base
RUN apk -U add curl
WORKDIR /usr/src/app
ENV CONTENT_API_URL http://localhost:3001
EXPOSE 3000

FROM node:dubnium AS build
RUN npm install -g @angular/cli@~8.3.4

WORKDIR /usr/src/app

# Install app dependencies
COPY package.json /usr/src/app/
RUN npm install

# Bundle app source
COPY . /usr/src/app
RUN ng build

FROM base as final
WORKDIR /usr/src/app
COPY --from=build /usr/src/app .
CMD [ "node", "app.js" ]
~
```

5. Press the Escape key and type `:wq` and then press the Enter key to save and close the file.

```
<Esc>
:wq
<Enter>
```

6. Rebuild the web application Docker image using the same command as you did previously.

```
docker image build -t content-web .
```

7. Create and start the image passing the correct URI to the API container as an environment variable. This variable will address the API application using its container name over the Docker network you created. After running the container, check to see the container is running and note the dynamic port assignment for the next step.

```
docker container run --name web --net fabmedical -P -d -e CONTENT_API_URL=http://api:3001 content-web
docker container ls
```

8. Curl the speakers path again, using the port assigned to the web container. Again, you will see HTML returned, but because curl does not process javascript, you cannot determine if the web application is communicating with the api application. You must verify this connection in a browser.

```
curl http://localhost:[PORT]/speakers.html
```

9. You will not be able to browse to the web application on the ephemeral port because the VM only exposes a limited port range. Now you will stop the web container and restart it using port 3000 to test in the browser. Type the following commands to stop the container, remove it, and run it again using explicit settings for the port.

```
docker container stop web
docker container rm web
docker container run --name web --net fabmedical -p 3000:3000 -d -e CONTENT_API_URL=http://api:3001 content-web
```

10. Curl the speaker path again, using port 3000. You will see the same HTML returned.

```
curl http://localhost:3000/speakers.html
```

11. You can now use a web browser to navigate to the website and successfully view the application at port 3000. Replace [BUILDAgentIP] with the IP address you used previously.

http://[BUILDAgentIP]:3000

EXAMPLE: http://13.68.113.176:3000

12. Commit your changes and push to the repository.

```
git add .
git commit -m "Setup Environment Variables"
git push
```

Enter credentials if prompted.

Task 6: Run several containers with Docker compose

Managing several containers with all their command line options can become difficult as the solution grows. `docker-compose` allows us to declare options for several containers and run them together.

1. First, cleanup the existing containers.

```
docker container stop web && docker container rm web
docker container stop api && docker container rm api
docker container stop mongo && docker container rm mongo
```

2. Navigate to your home directory (where you checked out the content repositories) and create a docker compose file.

```
cd ~
vi docker-compose.yml
<i>
```

Type the following as the contents of `docker-compose.yml`:

```
version: "3.4"

services:
  mongo:
    image: mongo
    restart: always

  api:
    build: ./content-api
    image: content-api
    depends_on:
      - mongo
    environment:
      MONGODB_CONNECTION: mongodb://mongo:27017/contentdb

  web:
    build: ./content-web
    image: content-web
    depends_on:
      - api
    environment:
      CONTENT_API_URL: http://api:3001
    ports:
      - "3000:3000"
```

Press the Escape key and type `:wq` and then press the Enter key to save and close the file.

```
<Esc>
:wq
<Enter>
```

3. Start the applications with the `up` command.

```
docker-compose -f docker-compose.yml -p fabmedical up -d
```

```
admin@fabmedical-[REDACTED]:~$ docker-compose -f docker-compose.yml -p fabmedical up -d
Creating network "fabmedical_default" with the default driver
Creating fabmedical_mongo_1 ... done
Creating fabmedical_api_1   ... done
Creating fabmedical_web_1  ... done
admin@fabmedical-[REDACTED]:~$
```

4. Visit the website in the browser; notice that we no longer have any data on the speakers or sessions pages.



speakers

5. We stopped and removed our previous mongodb container; all the data contained in it has been removed. Docker compose has created a new, empty mongodb instance that must be reinitialized. If we care to persist our data between container instances, docker has several mechanisms to do so. First, we will update our compose file to persist mongodb data to a directory on the build agent.

```
mkdir data
vi docker-compose.yml
```

Update the mongo service to mount the local data directory onto to the `/data/db` volume in the docker container.

```
mongo:
  image: mongo
  restart: always
  volumes:
    - ./data:/data/db
```

The result should look similar to the following screenshot:

vi docker-compose.init.yml

Add the following as the content:

```
version: "3.4"

services:
  init:
    build: ./content-init
    image: content-init
    depends_on:
      - mongo
  environment:
    MONGODB CONNECTION: mongodb://mongo:27017/contentdb
```

7. To reconfigure the mongodb volume, we need to bring down the mongodb service first.

```
docker-compose -f docker-compose.yml -p fabmedical down
```

```
admin@fabmedical-fabmedical-:~$ docker-compose -f docker-compose.yml -p fabmedical down
Stopping fabmedical_web_1 ... done
Stopping fabmedical_api_1 ... done
Stopping fabmedical_mongo_1 ... done
Removing fabmedical_web_1 ... done
Removing fabmedical_api_1 ... done
Removing fabmedical_mongo_1 ... done
Removing network fabmedical_default
admin@fabmedical-fabmedical-:~$
```

8. Now run `up` again with both files to update the mongodb configuration and run the initialization script.

```
docker-compose -f docker-compose.yml -f docker-compose.init.yml -p fabmedical up -d
```

9. Check the data folder to see that mongodb is now writing data files to the host.

```
ls ./data/
```

```
admin@fabmedical-fabmedical-sol:~$ docker-compose -f docker-compose.yml -f docker-compose.init.yml -p fabmedical up -d
Creating network "fabmedical_default" with the default driver
Creating fabmedical_mongo_1 ... done
Creating fabmedical_api_1 ... done
Recreating fabmedical_init_1 ... done
Creating fabmedical_web_1 ... done
admin@fabmedical-fabmedical-sol:~$ ls ./data/
collection-0--7350346735947181636.wt index-1--7350346735947181636.wt _mdb_catalog.wt WiredTigerLAS.wt
collection-2--7350346735947181636.wt index-3--7350346735947181636.wt mongod.lock WiredTiger.lock
collection-4--7350346735947181636.wt index-5--7350346735947181636.wt sizeStorer.wt WiredTiger.turtle
collection-6--7350346735947181636.wt index-7--7350346735947181636.wt storage.bson WiredTiger.wt
diagnostic.data journal WiredTiger
admin@fabmedical-fabmedical-sol:~$
```

10. Check the results in the browser. The speaker and session data are now available.

CONTOSO NEURO 2017

[Speakers](#) [Sessions](#)

SEPTEMBER 14-17, 2017

Monterey Conference Center
Monterey, California

 sessions

 Improves Motor Function and CNS Biomarkers in PD: Results from a Phase 2A Pilot Trial

Theresa Zesiewicz - Kevin Allison - Israt Jahan - Jessica Shaw - F. Reed Murtagh - Tracy Jones - Clifton Gooch - Jason Salemi - Matthew B Klein - Guy Miller - Kelly Sullivan

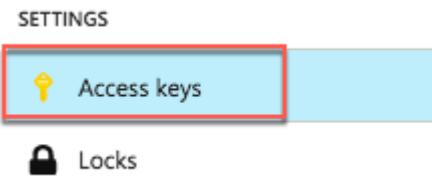
Task 7: Push images to Azure Container Registry

To run containers in a remote environment, you will typically push images to a Docker registry, where you can store and distribute images. Each service will have a repository that can be pushed to and pulled from with Docker commands. Azure Container Registry (ACR) is a managed private Docker registry service based on Docker Registry v2.

In this task, you will push images to your ACR account, version images with tagging, and setup continuous integration (CI) to build future versions of your containers and push them to ACR automatically.

1. In the [Azure Portal](#), navigate to the ACR you created in Before the hands-on lab.

2. Select **Access keys** under **Settings** on the left-hand menu.



3. The Access keys blade displays the Login server, username, and password that will be required for the next step. Keep this handy as you perform actions on the build VM.

Note: If the username and password do not appear, select Enable on the Admin user option.

4. From the cloud shell session connected to your build VM, login to your ACR account by typing the following command. Follow the instructions to complete the login.

```
docker login [LOGINSERVER] -u [USERNAME] -p [PASSWORD]
```

For example:

```
docker login fabmedical.soll.azurecr.io -u fabmedical -p +W/j=1+Fcze=n07SchxvGS1vsLRh/7ga
```

A screenshot of a terminal window titled 'Bash'. The user is running the command 'docker login fabmedical.soll.azurecr.io -u fabmedical -p +W/j=1+Fcze=n07SchxvGS1vsLRh/7ga'. The output shows several warning messages: 'WARNING! Using --password via the CLI is insecure. Use --password-stdin.', 'WARNING! Your password will be stored unencrypted in /home/adminfabmedical/.docker/config.json.', and 'Configure a credential helper to remove this warning. See https://docs.docker.com/engine/reference/commandline/login/#credentials-store'. Finally, it shows 'Login Succeeded'.

Tip: Make sure to specify the fully qualified registry login server (all lowercase).

5. Run the following commands to properly tag your images to match your ACR account name.

```
docker image tag content-web [LOGINSERVER]/content-web
docker image tag content-api [LOGINSERVER]/content-api
docker image tag content-init [LOGINSERVER]/content-init
```

6. List your docker images and look at the repository and tag. Note that the repository is prefixed with your ACR login server name, such as the sample shown in the screenshot below.

```
docker image ls
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
content-init	latest	48cf42890726	3 minutes ago	114MB
<none>	<none>	b8030e52a377	3 minutes ago	668MB
fabmedical.azurecr.io/content-web	latest	e8e3435f523b	20 minutes ago	367MB
content-web	latest	e8e3435f523b	20 minutes ago	367MB
<none>	<none>	e07b9746e1c8	21 minutes ago	1.32GB
<none>	<none>	968deadc9984	47 minutes ago	367MB
<none>	<none>	b3c206b3d87a	49 minutes ago	1.32GB
content-api	latest	f18052e28158	About an hour ago	117MB
fabmedical.azurecr.io/content-api	latest	f18052e28158	About an hour ago	117MB
<none>	<none>	8eefc4be1141	About an hour ago	671MB
node	alpine	fac3d6a8e034	8 days ago	106MB
node	dubnium	d5680e53a228	11 days ago	903MB
node	dubnium-alpine	a0708430821e	2 weeks ago	75.4MB
mongo	latest	965553e202a4	4 weeks ago	363MB
node	argon	ef4b194d8fcf	19 months ago	653MB

7. Push the images to your ACR account with the following command:

```
docker image push [LOGINSERVER]/content-web  
docker image push [LOGINSERVER]/content-api  
docker image push [LOGINSERVER]/content-init
```

In this screenshot of the console window, an example of images being pushed to an ACR account results from typing and running the following at the command prompt: docker push [LOGINSERVER]/content-web.

8. In the Azure Portal, navigate to your ACR account, and select **Repositories** under **Services** on the left-hand menu. You will now see two, one for each image.

Refresh

Search to filter repositories ...

Repositories ↑↓

content-api

content-init

content-web

9. Select `content-api`. You will see the latest tag is assigned.

ositories

«

⟳ Refresh

🔍 Search to filter repositories ...

Repositories ↑↓

content-api

...

content-init

...

content-web

...

content-api

Repository

⟳ Refresh

Delete

Repository

: content-api

Last updated date : 7/12/2020, 2:02 PM CDT

🔍 Search to filter tags ...

Tags ↑↓

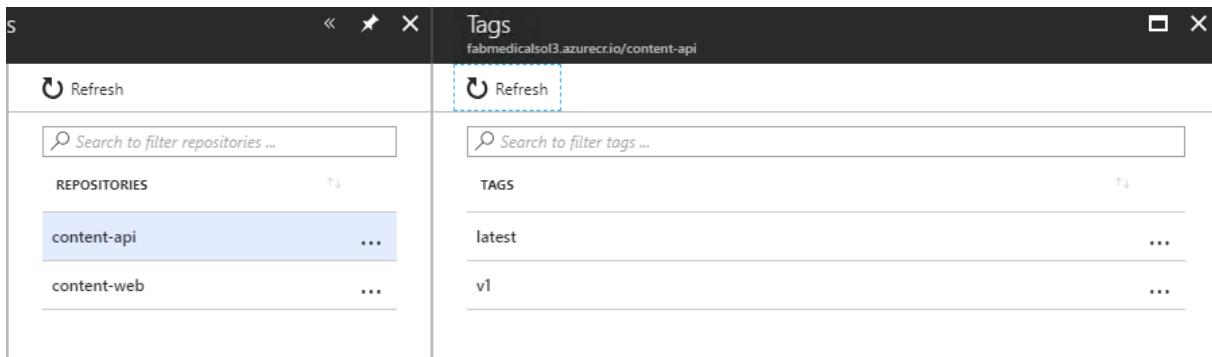
latest

10. From the cloud shell session attached to the VM, assign the `v1` tag to each image with the following commands. Then list the Docker images to note that there are now two entries for each image: showing the `latest` tag and the `v1` tag. Also note that the image ID is the same for the two entries, as there is only one copy of the image.

```
docker image tag [LOGINSERVER]/content-web:latest [LOGINSERVER]/content-web:v1
docker image tag [LOGINSERVER]/content-api:latest [LOGINSERVER]/content-api:v1
docker image ls
```

```
admin@fabmedical-fabmedical-cp4:~$ docker image ls
REPOSITORY          TAG      IMAGE ID      CREATED       SIZE
content-init        latest   4418620e6a4e  8 minutes ago  125MB
fabmedicalcp4.azurecr.io/content-init    latest   4418620e6a4e  8 minutes ago  125MB
fabmedicalcp4.azurecr.io/content-init    v1      4418620e6a4e  8 minutes ago  125MB
<none>              <none>   127d2e1f0417  8 minutes ago  671MB
content-web         latest   f8a5a30d7e01  25 minutes ago  354MB
fabmedicalcp4.azurecr.io/content-web    latest   f8a5a30d7e01  25 minutes ago  354MB
fabmedicalcp4.azurecr.io/content-web    v1      f8a5a30d7e01  25 minutes ago  354MB
<none>              <none>   1d065e18aaef  26 minutes ago  1.31GB
<none>              <none>   6e333b5cd0ae  35 minutes ago  354MB
<none>              <none>   440be66d5b98  36 minutes ago  1.31GB
content-api          latest   798d89b8f0aa  44 minutes ago  129MB
fabmedicalcp4.azurecr.io/content-api    latest   798d89b8f0aa  44 minutes ago  129MB
fabmedicalcp4.azurecr.io/content-api    v1      798d89b8f0aa  44 minutes ago  129MB
<none>              <none>   8a86289adce3  45 minutes ago  674MB
mongo               latest   409c3f937574  4 days ago    493MB
node                alpine   0f2c18cef5d3  11 days ago   117MB
node                dubnium  5ebbf4bb3837  2 weeks ago   911MB
node                dubnium-alpine  8e473595b853  4 weeks ago   83.5MB
node                argon    ef4b194d8fcf  2 years ago   653MB
admin@fabmedical-fabmedical-cp4:~$
```

11. Repeat Step 7 to push the images to ACR again so that the newly tagged `v1` images are pushed. Then refresh one of the repositories to see the two versions of the image now appear.



12. Run the following commands to pull an image from the repository. Note that the default behavior is to pull images tagged with `latest`. You can pull a specific version using the version tag. Also, note that since the images already exist on the build agent, nothing is downloaded.

```
docker image pull [LOGINSERVER]/content-web
docker image pull [LOGINSERVER]/content-web:v1
```

Task 8: Setup CI Pipeline to Push Images

In this task, you will use YAML to define a GitHub Actions workflow that builds your Docker image and pushes it to your ACR instance automatically.

1. In GitHub, return to the **Fabmedical** repository screen, and select the **Settings** tab.
2. From the left menu, select **Secrets**.
3. Select the **New secret** button.

The screenshot shows the GitHub interface for the 'Fabmedical' repository. The top navigation bar includes options like 'Code', 'Issues', 'Pull requests', 'Actions', 'Projects', 'Wiki', 'Security', 'Insights', and 'Settings'. The 'Settings' tab is highlighted with a red box. Below the navigation, there's a sidebar with links: 'Options', 'Manage access', 'Security & analysis', 'Branches', 'Webhooks', 'Notifications', 'Integrations', 'Deploy keys', 'Autolink references', 'Secrets' (which is also highlighted with a red box), and 'Actions'. The main content area is titled 'Secrets' and contains the message: 'There are no secrets for this repository.' It also notes that 'Encrypted secrets allow you to store sensitive information, such as access tokens, in your repository.' A 'New secret' button is located in the top right corner of the main content area, also highlighted with a red box.

4. In the **New secret** form, enter the name `ACR_USERNAME` and for the value, paste in the Azure Container Registry **Username** that was copied previously. Select **Add secret**.

Secrets / New secret

Name

Value

Add secret

5. Add another Secret, by entering the name `ACR_PASSWORD` and for the value, paste in the Azure Container Registry **Password** that was copied previously.

Secrets

[New secret](#)

Secrets are environment variables that are **encrypted** and only exposed to selected actions. Anyone with **collaborator** access to this repository can use these secrets in a workflow.

Secrets are not passed to workflows that are triggered by a pull request from a fork. [Learn more](#).

ACR_PASSWORD	Updated now	Update	Remove
ACR_USERNAME	Updated 23 seconds ago	Update	Remove

6. In your Azure Cloud Shell session connected to the build agent VM, navigate to the `~/Fabmedical` directory:

```
cd ~/Fabmedical
```

7. Before the GitHub Actions workflows can be setup, the `.github/workflows` directory needs to be created. Do this by running the following commands:

```
mkdir ~/Fabmedical/.github  
mkdir ~/Fabmedical/.github/workflows
```

8. Navigate to the `.github/workflows` directory:

```
cd ~/Fabmedical/.github/workflows
```

9. Next create the workflow YAML file.

```
vi content-web.yml
```

Add the following as the content. Be sure to replace the following placeholders:

- o replace [SHORT_SUFFIX] with your short suffix such as SOL .

```
name: content-web

# This workflow is triggered on push to the 'content-web' directory of the master branch of the repository
on:
  push:
    branches:
      - master
    paths:
      - 'content-web/**'

# Configure workflow to also support triggering manually
workflow_dispatch:
  inputs:
    logLevel:
      description: 'Log level'
      required: true
      default: 'warning'

# Environment variables are defined so that they can be used throughout the job definitions.
env:
  imageRepository: 'content-web'
  resourceGroupName: 'Fabmedical-[SHORT_SUFFIX]'
  containerRegistryName: 'fabmedical[SHORT_SUFFIX]'
  containerRegistry: 'fabmedical[SHORT_SUFFIX].azurecr.io'
  dockerfilePath: './content-web'
  tag: '${{ github.run_id }}'

# Jobs define the actions that take place when code is pushed to the master branch
jobs:
  build-and-publish-docker-image:
    name: Build and Push Docker Image
    runs-on: ubuntu-latest
    steps:
      # Checkout the repo
      - uses: actions/checkout@master

      - name: Build and push an image to container registry
        uses: docker/build-push-action@v1
        with:
          username: ${{ secrets.ACR_USERNAME }}
          password: ${{ secrets.ACR_PASSWORD }}
          path: ${{ env.dockerfilePath }}
          dockerfile: '${{ env.dockerfilePath }}/Dockerfile'
          registry: ${{ env.containerRegistry }}
          repository: ${{ env.imageRepository }}
          tags: ${{ env.tag }},latest</code></pre></li>
```

0. Save the file and exit VI by pressing <Esc> then :wq.

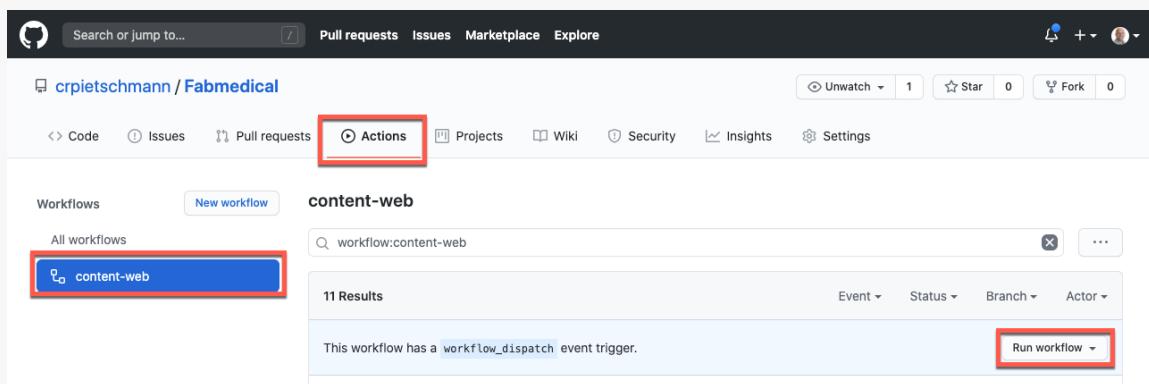
1. Save the pipeline YAML, then commit and push it to the Git repository:

```
git add .
git commit -m "Added workflow YAML"
git push
```

2. In GitHub, return to the `Fabmedical` repository screen, and select the **Actions** tab.

3. On the **Actions** page, select the `content-web` workflow.

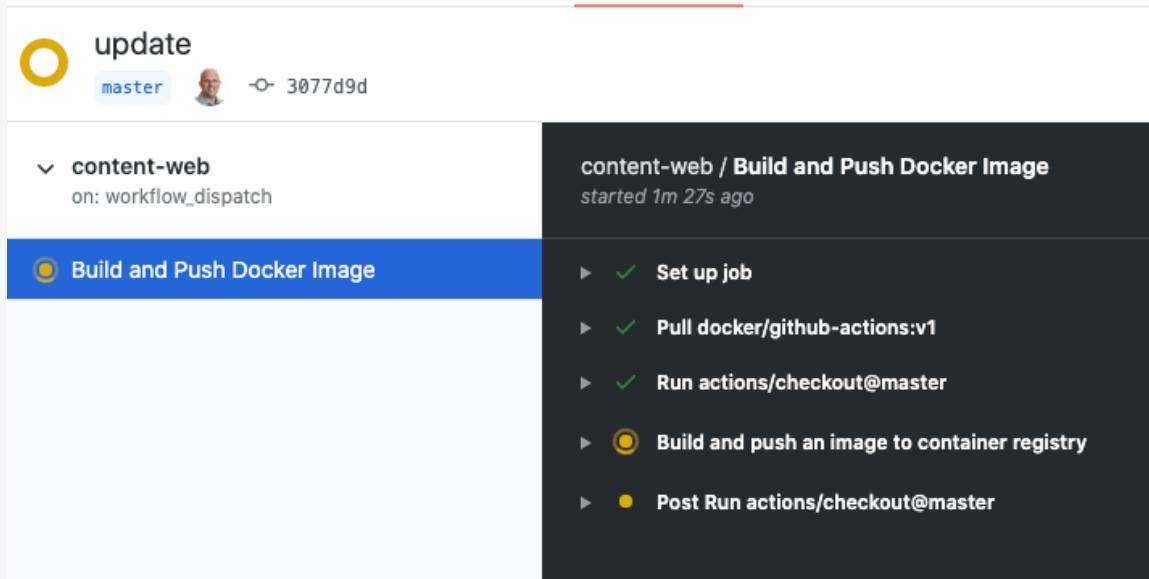
4. On the `content-web` workflow, select **Run workflow** and manually trigger the workflow to execute.



The screenshot shows the GitHub Actions interface for the repository `cpietschmann/Fabmedical`. The **Actions** tab is active. In the **content-web** workflow list, the first item, `content-web`, is highlighted with a red box. At the bottom right of the workflow card, there is a button labeled **Run workflow**, which is also highlighted with a red box.

5. After a second, the newly triggered workflow execution will display in the list. Select the new `content-web` execution.

6. Selecting the `Build and Push Docker Image` job of the workflow will display its execution status.



The screenshot shows the execution details for the `content-web / Build and Push Docker Image` job. The job was started 1m 27s ago. The steps listed are:

- Set up job
- Pull docker/github-actions:v1
- Run actions/checkout@master
- Build and push an image to container registry (status: in progress)
- Post Run actions/checkout@master

7. Next, setup the `content-api` workflow. This repository already includes `content-api` located within the `.github/workflows` directory. Open the file and review the code.

8. Edit the `resourceGroupName` and `containerRegistry` environment values to replace `[SHORT_SUFFIX]` with your own three-letter suffix.

```
    required: true
    default: 'warning'

# Environment variables are defined so that they can be used throughout the job definitions.
env:
  imageRepository: 'content-api'
  resourceGroupName: 'Fabmedical-[SHORT-SUFFIX]'
  containerRegistry: 'fabmedical[SHORT_SUFFIX].azurecr.io'
  dockerfilePath: './content-api'
  tag: '${{ github.run_id }}'

# Jobs define the actions that take place when code is pushed to the master branch
jobs:
```

9. Save the file, then navigate to the repositories in GitHub, select Actions, and then manually run the `content-api` workflow.

10. Next, setup the `content-init` workflow. Follow the same steps as the previous `content-api` workflow for the `content-init` workflow.

Exercise 2: Deploy the solution to Azure Kubernetes Service

Duration: 30 minutes

In this exercise, you will connect to the Azure Kubernetes Service cluster you created before the hands-on lab and deploy the Docker application to the cluster using Kubernetes.

Task 1: Tunnel into the Azure Kubernetes Service cluster

In this task, you will gather the information you need about your Azure Kubernetes Service cluster to connect to the cluster and execute commands to connect to the Kubernetes management dashboard from cloud shell.

Note: The following tasks should be executed in cloud shell and not the build machine, so disconnect from build machine if still connected.

1. Verify that you are connected to the correct subscription with the following command to show your default subscription:

```
az account show
```

- If you are not connected to the correct subscription, list your subscriptions and then set the subscription by its id with the following commands (similar to what you did in cloud shell before the lab):

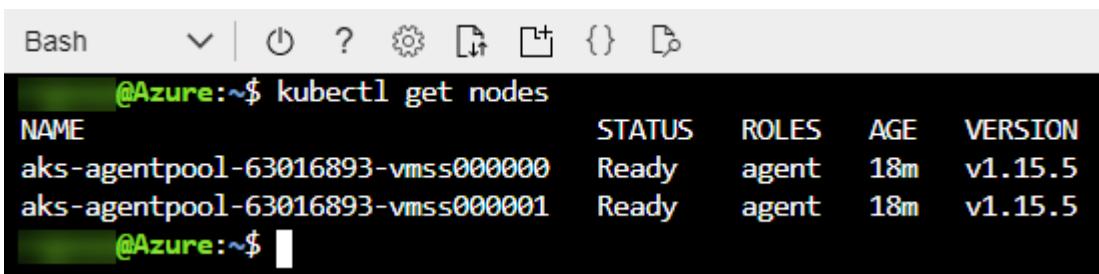
```
az account list
az account set --subscription {id}
```

2. Configure kubectl to connect to the Kubernetes cluster:

```
az aks get-credentials -a --name fabmedical-SUFFIX --resource-group fabmedical-SUFFIX
```

3. Test that the configuration is correct by running a simple kubectl command to produce a list of nodes:

```
kubectl get nodes
```



NAME	STATUS	ROLES	AGE	VERSION
aks-agentpool-63016893-vmss00000	Ready	agent	18m	v1.15.5
aks-agentpool-63016893-vmss00001	Ready	agent	18m	v1.15.5

4. Since the AKS cluster uses RBAC, a `ClusterRoleBinding` must be created before you can correctly access the dashboard.
To create the required binding, execute the command below:

```
kubectl create clusterrolebinding kubernetes-dashboard --clusterrole=cluster-admin --serviceaccount=kube-system:kube
```

Note: If you get an error saying `error: failed to create clusterrolebinding: clusterrolebindings.rbac.authorization.k8s.io "kubernetes-dashboard" already exists` just ignore it and move on to the next step.

5. Before you can create an SSH tunnel and connect to the Kubernetes Dashboard, you will need to download the **Kubeconfig** file within Azure Cloud Shell that contains the credentials you will need to authenticate to the Kubernetes Dashboard.

Within the Azure Cloud Shell, use the following command to download the Kubeconfig file:

```
download /home/<username>/.kube/config
```

Make sure to replace the `<username>` placeholder with your name from the command-line in the Azure Cloud Shell.

Note: You can find the `<username>` from the first part of the Azure Cloud Shell command-line prompt; such as `<username>@Azure:~$`.

You can also look in the `/home` directory and see the directory name that exists within it to find the correct username directory where the Kubeconfig file resides:

```
ls /home
```

6. Create an SSH tunnel linking a local port (`8001`) on your cloud shell host to port `443` on the management node of the cluster. Cloud shell will then use the web preview feature to give you remote access to the Kubernetes dashboard. Execute the command below replacing the values as follows:

Note: After you run this command, it may work at first and later lose its connection, so you may have to run this again to reestablish the connection. If the Kubernetes dashboard becomes unresponsive in the browser this is an indication to return here and check your tunnel or rerun the command.

```
az aks browse --name fabmedical-SUFFIX --resource-group fabmedical-SUFFIX
```

7. If the tunnel is successful, you will see the Kubernetes Dashboard authentication screen. Select the **Kubeconfig** option, select the ellipsis (...) button, select the **Kubeconfig** file that was previously downloaded, then select **Sign in**.

Kubernetes Dashboard

Token
Every Service Account has a Secret with valid Bearer Token that can be used to log in to Dashboard. To find out more about how to configure and use Bearer Tokens, please refer to the [Authentication](#) section.

Kubeconfig
Please select the kubeconfig file that you have created to configure access to the cluster. To find out more about how to configure and use kubeconfig file, please refer to the [Configure Access to Multiple Clusters](#) section.

Choose kubeconfig file ...

Sign in

8. Once authenticated, you will see the Kubernetes management dashboard.

The screenshot shows the Kubernetes Dashboard's Overview page. On the left, a sidebar lists navigation items: Cluster (Namespaces, Nodes, Persistent Volumes, Roles, Storage Classes), Namespace (default dropdown), and Workloads (Cron Jobs, Daemon Sets). The main content area has two sections: "Discovery and Load Balancing" (Services) and "Config and Storage" (Secrets).

Discovery and Load Balancing

Services

Name	Labels	Cluster IP	Internal endpoints	External endpoints	Age	⋮
kubernetes	component=... provider=...	10.0.0.1	kubernetes:4 kubernetes:0	-	14 hours	⋮

Config and Storage

Secrets

Name	Type	Age	⋮
default-token-s6kmco	kubernetes.io/service-account	14 hours	⋮

Note: If the tunnel is not successful (if a JSON output is displayed), execute the command below and then return to task 5 above:

```
az extension add --name aks-preview
```

Task 2: Deploy a service using the Kubernetes management dashboard

In this task, you will deploy the API application to the Azure Kubernetes Service cluster using the Kubernetes dashboard.

1. From the Kubernetes dashboard, select **Create** in the top right corner.

2. From the Resource creation view, select **Create an App**.

The screenshot shows the Kubernetes Management Dashboard's 'Create' interface for deploying an application. The 'Create from form' tab is active. The 'App name' field is set to 'api'. The 'Container image' field contains the URL 'fabmedicalcp3.azurecr.io/content-api'. The 'Number of pods' field is set to '1'. The 'Service' dropdown is set to 'Internal'. In the 'Advanced Options' section, the 'Port' and 'Target port' are both set to '3001' and 'Protocol' is set to 'TCP'. At the bottom, there are 'Deploy', 'Cancel', and 'Show advanced options' buttons.

- Enter `api` for the App name.
- Enter `[LOGINSERVER]/content-api` for the Container Image, replacing `[LOGINSERVER]` with your ACR login server, such as `fabmedicalsol.azurecr.io`.
- Set Number of pods to `1`.
- Set Service to `Internal`.
- Use `3001` for Port and `3001` for Target port.

3. Select **SHOW ADVANCED OPTIONS**

- Enter `1` for the CPU requirement.
- Enter `128` for the Memory requirement.

Description	The description will be added as an annotation to the Deployment and displayed in the application's details.						
Labels <table border="1"> <tr> <td>key</td> <td>value</td> </tr> <tr> <td>k8s-app</td> <td>api</td> </tr> </table> <p>3 / 253</p> <table border="1"> <tr> <td>key</td> <td>value</td> </tr> </table> <p>0 / 253</p>		key	value	k8s-app	api	key	value
key	value						
k8s-app	api						
key	value						
Namespace *	Namespaces let you partition resources into logically named groups. Learn more						
default	▼						
Image Pull Secret <div style="border: 1px solid #ccc; padding: 5px; margin-bottom: 5px;">▼</div>							
CPU requirement (cores)	Memory requirement (MiB)						
1	128						
Run command <div style="border: 1px solid #ccc; padding: 5px; margin-top: 5px;">▼</div>							

4. Select Deploy to initiate the service deployment based on the image. This can take a few minutes. In the meantime, you will be redirected to the Overview dashboard. Select the **API** deployment from the **Overview** dashboard to see the deployment in progress.

☰ Overview

Cluster

- Namespaces
- Nodes
- Persistent Volumes
- Roles
- Storage Classes

Namespace

default ▾

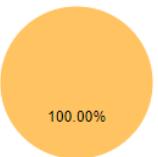
Overview

Workloads

- Cron Jobs
- Daemon Sets
- Deployments
- Jobs
- Pods
- Replica Sets
- Replication Controllers
- Stateful Sets

Workloads

Workloads Statuses



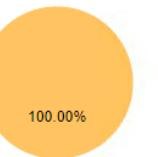
100.00%

Deployments



100.00%

Pods



100.00%

Replica Sets

Deployments

Name	Labels	Pods	Age	Images
api	k8s-app: api	0 / 1	0 seconds	fabmedicalsol.azurecr.io/

Pods

Name	Node	Status	Restarts	Age	CPU (cores)	Memory (bytes)
api-5dddfdf	aks-agentpool-30625191-1	Waiting: Con	0	0 seconds	-	-

5. Kubernetes indicates a problem with the `api` Replica Set after some seconds. Select the log icon to investigate.

☰ Workloads > Deployments > api

SCALE EDIT DELETE

Cluster

- Namespaces
- Nodes
- Persistent Volumes
- Roles
- Storage Classes

Namespace

default ▾

Overview

Workloads

- Cron Jobs
- Daemon Sets
- Deployments
- Jobs
- Pods
- Replica Sets
- Replication Controllers
- Stateful Sets

Details

Name: api
Namespace: default
Labels: k8s-app: api
Annotations: deployment.kubernetes.io/revision: 1
Creation Time: 2018-06-11T12:32 UTC
Selector: k8s-app: api
Strategy: RollingUpdate
Min ready seconds: 0
Revision history limit: 10
Rolling update strategy: Max surge: 25%, Max unavailable: 25%
Status: 1 updated, 1 total, 0 available, 1 unavailable

New Replica Set

Name	Labels	Pods	Age	Images
! api-67bbbfcd	k8s-app: api pod-template-hash: 2366.	1 / 1	39 seconds	fabmedicalsol3.azurecr.io/

Old Replica Sets

6. The log indicates that the content-api application is once again failing because it cannot find a mongodb instance to communicate with. You will resolve this issue by migrating your data workload to Cosmos DB.

Logs from api in api-7f45f6fd75-g92r2

```
> content-api@0.0.0 start /usr/src/app
> node ./server.js
Listening on port 3001
Could not connect to MongoDB!
MongoTimeoutError: Server selection timed out after 30000 ms
npm ERR! code ELIFECYCLE
npm ERR! errno 255
npm ERR! content-api@0.0.0 start: `node ./server.js`
npm ERR! Exit status 255
npm ERR!
npm ERR! Failed at the content-api@0.0.0 start script.
npm ERR! This is probably not a problem with npm. There is likely additional logging output above.
npm ERR! A complete log of this run can be found in:
npm ERR!     /root/.npm/_logs/2019-12-05T00_57_24_656Z-debug.log
```

7. Open the Azure portal in your browser and navigate to your resource group and find your Cosmos DB resource. Select the Cosmos DB resource to view details.

<input type="checkbox"/>	NAME	TYPE	LOCATION
<input type="checkbox"/>	fabmedical-sol	Virtual machine	East US
<input type="checkbox"/>	fabmedical-sol	Kubernetes service	East US
<input checked="" type="checkbox"/>	fabmedical-sol OSDISK_1_30214e494c704050bde59c1d45a9**	DISK	East US
<input type="checkbox"/>	fabmedical-sol2	Azure Cosmos DB account	East US
<input type="checkbox"/>	fabmedical-sol3	Container registry	East US
<input type="checkbox"/>	fabmedical-sol637	Network interface	East US
<input type="checkbox"/>	fabmedical-sol-ip	Public IP address	East US
<input type="checkbox"/>	fabmedical-sol-nsg	Network security group	East US
<input type="checkbox"/>	fabmedical-sol-vnet	Virtual network	East US

8. Under **Quick Start** select the **Node.js** tab and copy the Node 3.0 connection string.

fabmedical-sol2 - Quick start

Azure Cosmos DB account

Search (Ctrl+ /)

Overview

Activity log

Access control (IAM)

Tags

Diagnose and solve problems

Quick start

Data Explorer

SETTINGS

Connection String

Preview Features

Replicate data globally

Default consistency

Firewall and virtual networks

Locks

Automation script

COLLECTIONS

Browse

Scale

Congratulations! Your Azure Cosmos DB account with MongoDB API is ready.

Now, let's connect your existing MongoDB app to it:

Choose a platform

.NET Node.js MongoDB Shell Java Python Others

1 Using the Node.js 2.2 driver, connect your existing MondoDB app

You can use your existing MongoDB Nodejs 2.2 driver to work with Azure Cosmos DB. Make sure to enable SSL. Here is an example

```
var mongoClient = require("mongodb").MongoClient;
mongoClient.connect("mongodb://fabmedical-sol2:2nb7uXNPKeg5NT73Vo400sjVThUTxE1pggqkr5J7guvVJB5B
db.close();
});
```

PRIMARY CONNECTION STRING

```
mongodb://fabmedical-sol2:2nb7uXNPKeg5NT73Vo400sjVThUTxE1pggqkr5J7guvVJB5B1JDsLJBcdtir1LUJVzthJq...
```

For more details on configuring Nodejs driver to use SSL, follow [this article](#).

Questions? [Contact us](#)

Using the Node.js 3.0 driver, connect your existing MondoDB app

You can use your existing MongoDB Nodejs 3.0 driver to work with Azure Cosmos DB. Make sure to enable SSL. Here is an example

```
var mongoClient = require("mongodb").MongoClient;
mongoClient.connect("mongodb://fabmedical-sol2:2nb7uXNPKeg5NT73Vo400sjVThU
client.close();
});
```

PRIMARY CONNECTION STRING

```
mongodb://fabmedical-sol2:2nb7uXNPKeg5NT73Vo400sjVThUTxE1pggqkr5J7guvVJB5B1JDsLJBcdtir1LUJVzthJq...
```

For more details on configuring Nodejs driver to use SSL, follow [this article](#).

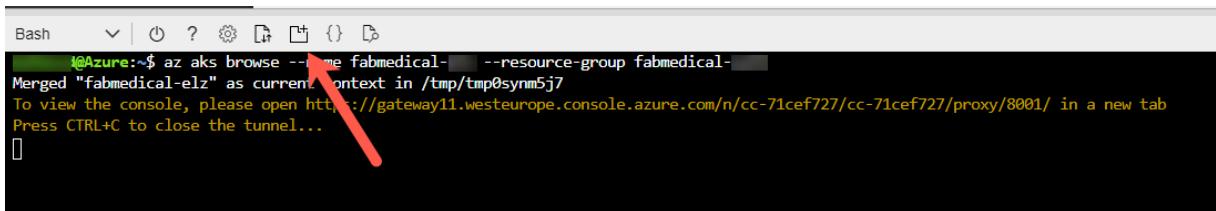
Questions? [Contact us](#)

9. Update the provided connection string with a database **contentdb** and a replica set **globaldb**.

Note: Username and password redacted for brevity.

`mongodb://<USERNAME>:<PASSWORD>@fabmedical-<SUFFIX>.documents.azure.com:10255/contentdb?ssl=true&replicaSet=globaldb`

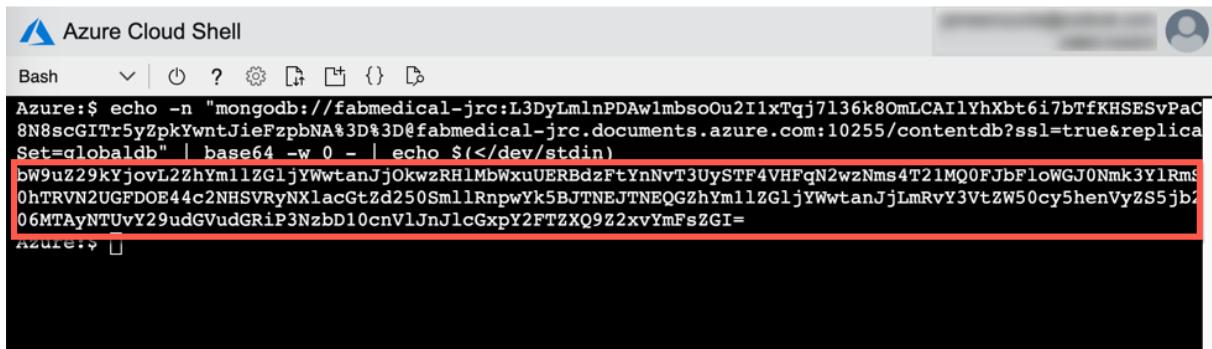
10. To avoid disconnecting from the Kubernetes dashboard, open a **new** Azure Cloud Shell console.



11. You will setup a Kubernetes secret to store the connection string and configure the content-api application to access the secret. First, you must base64 encode the secret value. Open your Azure Cloud Shell window and use the following command to encode the connection string and then, copy the output.

Note: Double quote marks surrounding the connection string are required to successfully produce the required output.

```
echo -n "[CONNECTION STRING VALUE]" | base64 -w 0 - | echo $(</dev/stdin)
```

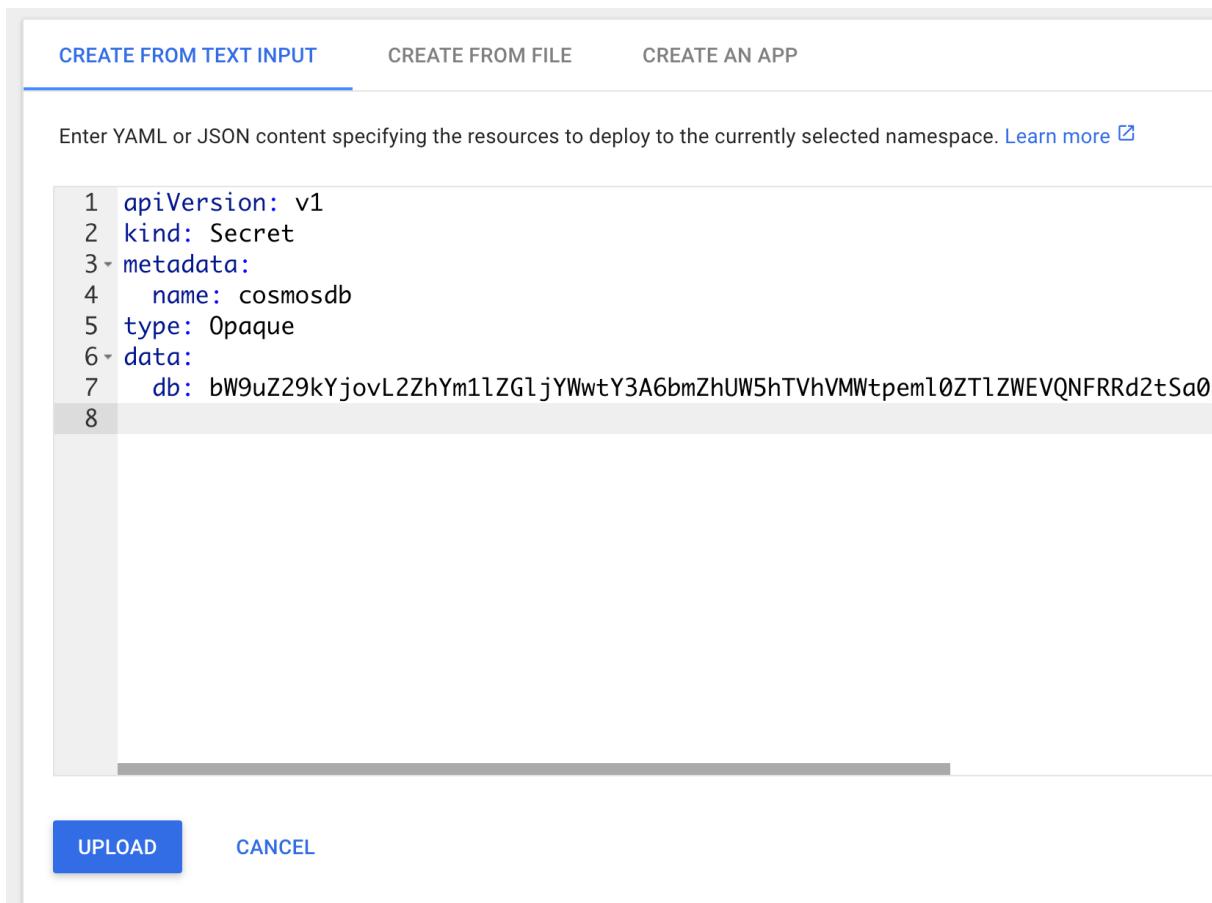


```
Azure:~$ echo -n "mongodb://fabmedical-jrc:L3DyLmlnPDAwlmbsoOu2I1xTqj7136k80mLCAI1YhXbt6i7bTfKHSEsvPaC8N8scGIt5yZpkYwntJieFzpbNA3D%3D&fabmedical-jrc.documents.azure.com:10255/contentdb?ssl=true&replicaSet=globaldb" | base64 -w 0 - | echo $(
```

The terminal window shows a command being run to echo a MongoDB connection string. The string contains sensitive information like a password and port, which is highlighted with a red box.

12. Return to the Kubernetes UI in your browser and select **+ Create**. Update the following YAML with the encoded connection string from your clipboard, paste the YAML data into the create dialog, and choose **Upload**.

```
apiVersion: v1
kind: Secret
metadata:
  name: cosmosdb
type: Opaque
data:
  db: <base64 encoded value>
```



The screenshot shows the 'CREATE FROM TEXT INPUT' tab selected in the Kubernetes UI. A large text area contains the YAML secret definition. Below the text area are two buttons: 'UPLOAD' (in blue) and 'CANCEL'.

```
1 apiVersion: v1
2 kind: Secret
3 metadata:
4   name: cosmosdb
5 type: Opaque
6 data:
7   db: bW9uZ29kYjovL2ZhYm1lZGljYWwtY3A6bmZhUW5hTVhVMWtpeml0ZTlZWEVQNFRd2tSa0f
8
```

13. Scroll down in the Kubernetes dashboard until you can see **Secrets** in the left-hand menu. Select it.

The screenshot shows the Kubernetes UI for managing secrets. On the left, there's a sidebar with navigation links like Workloads, Discovery and Load Balancing, Config and Storage, and Secrets (which is currently selected). The main area displays a table of secrets:

Name	Type
cosmosdb	Opaque
default-token-h25sn	kubernetes.io/service-account-token

14. View the details for the **cosmosdb** secret. Select the eyeball icon to show the secret.

The screenshot shows the detailed view for the **cosmosdb** secret. It includes the following information:

Field	Value
Name:	cosmosdb
Namespace:	default
Creation Time:	2020-06-27T19:18 UTC
Type:	Opaque

Below the details, there's a section for Data, which contains a MongoDB connection string:

```
db: mongodb://fabmedical-cp:nfaQnaMXU1kizite9YXEP4TQwkRkJTAmdnGMNpKupvO3v0ghXBanesxtAUjKi7BjPMDy67pUmtzzQ53mxWJiCA%3D%3D@fabmedical-cp.documents.azure.com:10255/contentdb?ssl=true&replicaSet=globaldb
```

15. Next, download the api deployment configuration using the following command in your Azure Cloud Shell window:

```
kubectl get -o=yaml deployment api > api.deployment.yml
```

16. Edit the downloaded file using cloud shell code editor:

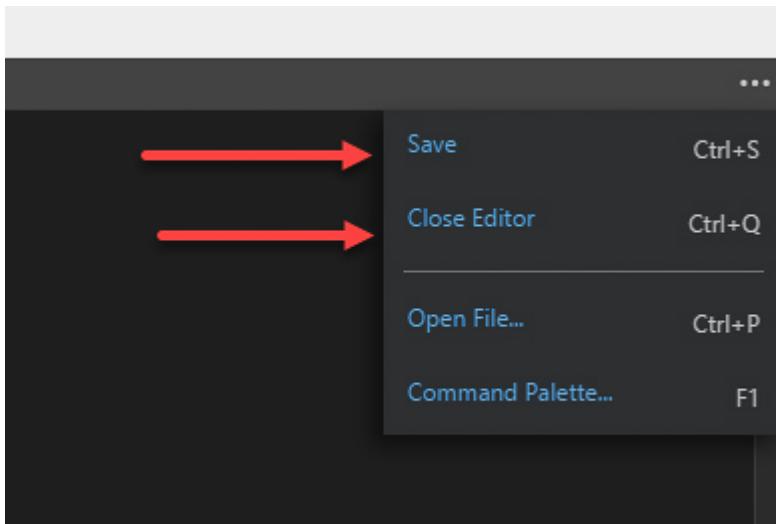
```
code api.deployment.yml
```

Add the following environment configuration to the container spec, below the "image" property:

```
env:
  - name: MONGODB_CONNECTION
    valueFrom:
      secretKeyRef:
        name: cosmosdb
        key: db

20  rollingupdate:
21    maxSurge: 25%
22    maxUnavailable: 25%
23    type: RollingUpdate
24  template:
25    metadata:
26      creationTimestamp: null
27      labels:
28        k8s-app: api
29      name: api
30  spec:
31    containers:
32      - image: fabmedicalcp.azurecr.io/content-api
33        env:
34          - name: MONGODB_CONNECTION
35            valueFrom:
36              secretKeyRef:
37                name: cosmosdb
38                key: db
39        imagePullPolicy: Always
40        name: api
41        resources:
42          requests:
43            cpu: 125m
44            memory: 128Mi
45        securityContext:
46          privileged: false
47        terminationMessagePath: /dev/termination-log
48        terminationMessagePolicy: File
49        dnsPolicy: ClusterFirst
```

17. Save your changes and close the editor.



18. Update the api deployment by using `kubectl` to apply the new configuration.

```
kubectl apply -f api.deployment.yml
```

19. Select **Deployments** then **api** to view the api deployment. It now has a healthy instance and the logs indicate it has connected to mongodb.

```
Logs from api      in api-7bb6777d69-bnlrv
&gt; content-api@0.0.0 start /usr/src/app
&gt; node ./server.js
Listening on port 3001
Connected to MongoDB
```

Task 3: Deploy a service using kubectl

In this task, deploy the web service using `kubectl`.

1. Open a **new** Azure Cloud Shell console.
2. Create a text file called `web.deployment.yml` using the Azure Cloud Shell Editor.

```
code web.deployment.yml
```

3. Copy and paste the following text into the editor:

Note: Be sure to copy and paste only the contents of the code block carefully to avoid introducing any special characters.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app: web
  name: web
```

```

spec:
  replicas: 1
  selector:
    matchLabels:
      app: web
  strategy:
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 1
    type: RollingUpdate
  template:
    metadata:
      labels:
        app: web
        name: web
    spec:
      containers:
        - image: [LOGINSERVER].azurecr.io/content-web
          env:
            - name: CONTENT_API_URL
              value: http://api:3001
          livenessProbe:
            httpGet:
              path: /
              port: 3000
            initialDelaySeconds: 30
            periodSeconds: 20
            timeoutSeconds: 10
            failureThreshold: 3
          imagePullPolicy: Always
          name: web
          ports:
            - containerPort: 3000
              hostPort: 80
              protocol: TCP
          resources:
            requests:
              cpu: 1000m
              memory: 128Mi
          securityContext:
            privileged: false
          terminationMessagePath: /dev/termination-log
          terminationMessagePolicy: File
      dnsPolicy: ClusterFirst
      restartPolicy: Always
      schedulerName: default-scheduler
      securityContext: {}
      terminationGracePeriodSeconds: 30

```

4. Update the `[LOGINSERVER]` entry to match the name of your ACR Login Server.

5. Select the ... button and choose **Save**.



6. Select the ... button again and choose **Close Editor**.



7. Create a text file called `web.service.yml` using the Azure Cloud Shell Editor.

```
code web.service.yml
```

8. Copy and paste the following text into the editor:

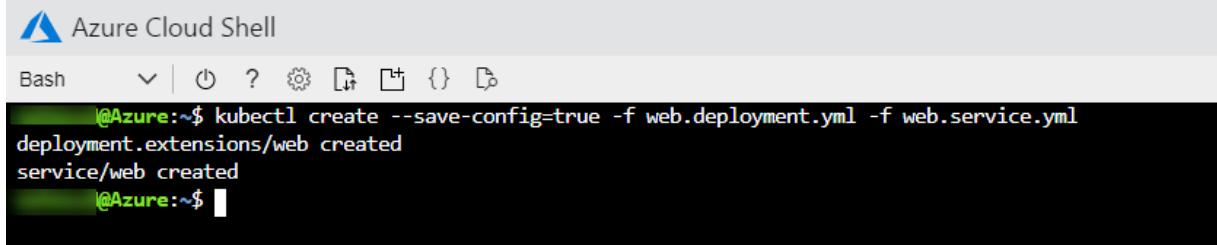
Note: Be sure to copy and paste only the contents of the code block carefully to avoid introducing any special characters.

```
apiVersion: v1
kind: Service
metadata:
  labels:
    app: web
    name: web
spec:
  ports:
    - name: web-traffic
      port: 80
      protocol: TCP
      targetPort: 3000
  selector:
    app: web
  sessionAffinity: None
  type: LoadBalancer
```

9. Save changes and close the editor.

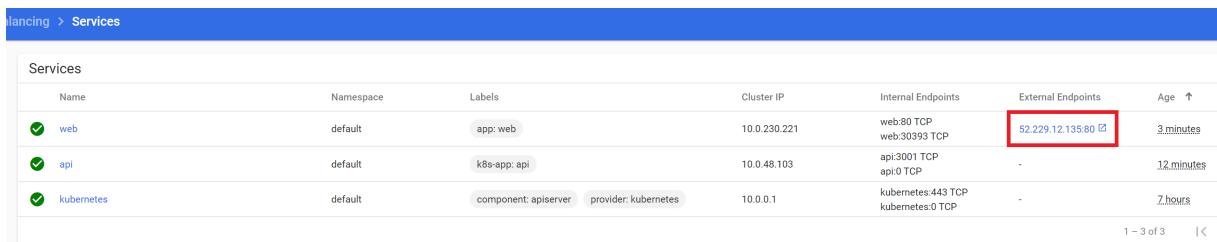
10. Type the following command to deploy the application described by the YAML files. You will receive a message indicating the items kubectl has created a web deployment and a web service.

```
kubectl create --save-config=true -f web.deployment.yml -f web.service.yml
```



Azure Cloud Shell interface showing a terminal window. The command `kubectl create --save-config=true -f web.deployment.yml -f web.service.yml` is run, resulting in the output: `deployment.extensions/web created` and `service/web created`. The terminal prompt is `@Azure:~$`.

11. Return to the browser where you have the Kubernetes management dashboard open. From the navigation menu, select Services view under Discovery and Load Balancing. From the Services view, select the web service, and from this view, you will see the web service deploying. This deployment can take a few minutes. When it completes, you should be able to access the website via an external endpoint.



The screenshot shows the Kubernetes Management Dashboard with the "Services" tab selected. A table lists three services: "web", "api", and "kubernetes". The "web" service is highlighted with a red box around its "External Endpoints" column, which shows the IP address `52.229.12.135:80`. The "api" service has no external endpoints listed. The "kubernetes" service also has no external endpoints listed. The table includes columns for Name, Namespace, Labels, Cluster IP, Internal Endpoints, External Endpoints, and Age.

Name	Namespace	Labels	Cluster IP	Internal Endpoints	External Endpoints	Age
web	default	app: web	10.0.230.221	web:80 TCP web:30393 TCP	52.229.12.135:80	3 minutes
api	default	k8s-app: api	10.0.48.103	api:3001 TCP api:0 TCP	-	12 minutes
kubernetes	default	component: apiserver provider: kubernetes	10.0.0.1	kubernetes:443 TCP kubernetes:0 TCP	-	7 hours

12. Select the speakers and sessions links. Note that no data is displayed, although we have connected to our Cosmos DB instance, there is no data loaded. You will resolve this by running the content-init application as a Kubernetes Job in Task 5.



sessions

Task 4: Deploy a service using a Helm chart

In this task, you will deploy the web service using a [Helm](#) chart to streamline the installing and managing the container-based application on the Azure Kubernetes cluster.

1. From the Kubernetes dashboard, under **Workloads**, select **Deployments**.
2. Select the triple vertical dots on the right of the `web` deployment and then choose **Delete**. When prompted, select **Delete** again.

Name	Labels	Pods	Age	Images
<code>web</code>	<code>app: web</code>	1 / 1	2 minutes	[redacted]
<code>api</code>	<code>k8s-app: api</code>	1 / 1	2 days	[redacted]

3. From the Kubernetes dashboard, under **Discovery and Load Balancing**, select **Services**.
4. Select the triple vertical dots on the right of the `web` service and then choose **Delete**. When prompted, select **Delete** again.

Name	Labels	Cluster IP	Internal endpoints	External endpoints	Age
web	app: web	10.0.137.94	web:80 TCP web:32365 TCP	23.101.136.188:... -	11 minutes
api	k8s-app: api	10.0.5.89	api:3001 TCP	-	2 days
kubernetes	component: apis. provider: kubern...	10.0.0.1	kubernetes:443... -	-	2 days

5. Open a **new** Azure Cloud Shell console.

6. Update your starter files by pulling the latest changes from the Git repository:

```
cd ~/MCW-Cloud-native-applications/Hands-on\ lab/lab-files/infrastructure/content-web
git pull
```

7. We will use the chart scaffold implementation that we have available in the source code. Use the following commands to access the chart folder:

```
cd ~/MCW-Cloud-native-applications/Hands-on\ lab/lab-files/infrastructure/content-web/charts
```

8. We now need to update the generated scaffold to match our requirements. We will first update the file named `values.yaml`.

```
cd web
code values.yaml
```

9. Search for the `image` definition and update the values so that they match the following:

```
image:
  repository: [LOGINSERVER].azurecr.io/content-web
  pullPolicy: Always
```

10. Search for `nameOverride` and `fullnameOverride` entries and update the values so that they match the following:

```
nameOverride: "web"
fullnameOverride: "web"
```

11. Search for the `service` definition and update the values so that they match the following:

```
service:
  type: LoadBalancer
```

```
port: 80
```

12. Search for the `resources` definition and update the values so that they match the following:

```
resources:  
  # We usually recommend not to specify default resources and to leave this as a conscious  
  # choice for the user. This also increases chances charts run on environments with little  
  # resources, such as Minikube. If you do want to specify resources, uncomment the following  
  # lines, adjust them as necessary, and remove the curly braces after 'resources:'.  
  # limits:  
  #   cpu: 100m  
  #   memory: 128Mi  
  requests:  
    cpu: 1000m  
    memory: 128Mi
```

13. Save changes and close the editor.

14. We will now update the file named `Chart.yaml`.

```
code Chart.yaml
```

15. Search for the `appVersion` entry and update the value so that it matches the following:

```
appVersion: latest
```

16. We will now update the file named `deployment.yaml`.

```
cd templates  
code deployment.yaml
```

17. Search for the `metadata` definition and update the values so that they match the following:

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  (...)  
spec:  
  (...)  
  template:  
    metadata:  
      (...)  
      annotations:  
        rollme: {{ randAlphaNum 5 | quote }}
```

18. Search for the `containers` definition and update the values so that they match the following:

```
containers:  
  - name: {{ .Chart.Name }}  
    securityContext:  
      {{- toYaml .Values.securityContext | nindent 12 }}  
    image: "{{ .Values.image.repository }}:{{ .Chart.AppVersion }}"  
    imagePullPolicy: {{ .Values.image.pullPolicy }}  
    ports:
```

```
- name: http
  containerPort: 3000
  protocol: TCP
env:
- name: CONTENT_API_URL
  value: http://api:3001
livenessProbe:
  httpGet:
    path: /
    port: 3000
```

19. Save changes and close the editor.

20. We will now update the file named `service.yaml`.

code service.yaml

21. Search for the `ports` definition and update the values so that they match the following:

```
ports:
- port: {{ .Values.service.port }}
  targetPort: 3000
  protocol: TCP
  name: http
```

22. Save changes and close the editor.

23. The chart is now setup to run our web container. Type the following command to deploy the application described by the YAML files. You will receive a message indicating that helm has created a web deployment and a web service.

```
cd ../..  
helm install web ./web
```

```
Bash └─? {} ⌂ @Azure:~/MCW-Cloud-native-applications/Hands-on lab/lab-files/developer/content-web/charts$ helm install web ./web
NAME: web
LAST DEPLOYED: Thu Dec 5 02:07:34 2019
NAMESPACE: default
STATUS: deployed
REVISION: 1
NOTES:
1. Get the application URL by running these commands:
   NOTE: It may take a few minutes for the LoadBalancer IP to be available.
   You can watch the status of by running 'kubectl get --namespace default svc -w web'
   export SERVICE_IP=$(kubectl get svc --namespace default web --template "{{ range (index .status.loadBalancer.ingress 0) }}{{.}}{{ end }}")
   echo http://$SERVICE_IP:80
@Azure:~/MCW-Cloud-native-applications/Hands-on lab/lab-files/developer/content-web/charts$
```

24. Return to the browser where you have the Kubernetes management dashboard open. From the navigation menu, select **Services** view under **Discovery and Load Balancing**. From the Services view, select the web service, and from this view, you will see the web service deploying. This deployment can take a few minutes. When it completes, you should be able to access the website via an external endpoint.

Services							
Name	Namespace	Labels	Cluster IP	Internal Endpoints	External Endpoints	Age	↑
 web	default	app: web	10.0.230.221	web:80 TCP web:30933 TCP	52.229.12.135:80	3 minutes	
 api	default	k8s-app: api	10.0.48.103	api:3001 TCP api:0 TCP	-	12 minutes	
 kubernetes	default	component: apiserver provider: kubernetes	10.0.0.1	kubernetes:443 TCP kubernetes:0 TCP	-	7 hours	

25. Select the speakers and sessions links. Note that no data is displayed, although we have connected to our Cosmos DB instance, there is no data loaded. You will resolve this by running the content-init application as a Kubernetes Job.



26. We will now persist the changes into the repository. Execute the following commands:

```
cd ..
git pull
git add charts/
git commit -m "Helm chart update."
git push
```

Task 5: Initialize database with a Kubernetes Job

In this task, you will use a Kubernetes Job to run a container that is meant to execute a task and terminate, rather than run all the time.

1. Create a text file called `init.job.yml` using the Azure Cloud Shell Editor.

```
code init.job.yml
```

2. Copy and paste the following text into the editor:

Note: Be sure to copy and paste only the contents of the code block carefully to avoid introducing any special characters.

```
apiVersion: batch/v1
kind: Job
metadata:
  name: init
spec:
  template:
    spec:
      containers:
        - name: init
          image: [LOGINSERVER]/content-init
          env:
            - name: MONGODB_CONNECTION
              valueFrom:
```

```
secretKeyRef:  
  name: cosmosdb  
  key: db  
restartPolicy: Never  
backoffLimit: 4
```

3. Edit this file and update the `[LOGINSERVER]` entry to match the name of your ACR Login Server.

4. Save changes and close the editor.

5. Type the following command to deploy the job described by the YAML. You will receive a message indicating the `kubectl` has created an init `job.batch`.

```
kubectl create --save-config=true -f init.job.yml
```

6. View the Job by selecting **Jobs** under **Workloads** in the Kubernetes UI.

The screenshot shows the Kubernetes UI with the navigation bar "Workloads > Jobs". On the left, there's a sidebar with "Cluster" and "Namespaces" sections, and a "Namespace" dropdown set to "default". Below that are sections for "Overview", "Workloads" (with "Cron Jobs", "Daemon Sets", "Deployments", and "Jobs" listed), and a "Logs" section. The main area is titled "Jobs" and lists one item:

Name	Labels	Pods	Age	Images
init	controller-uid: 2941c9f2-... job-name: init	0 / 1	8 seconds	fabmedicalsol3.azurecr.io/...

A red arrow points to the log icon (three vertical dots) next to the "init" job entry.

7. Select the log icon to view the logs.

Logs from init in init-n8k62

```
&gt; content-init@1.0.0 start /usr/src/app
&gt; node server.js
Clean Sessions table
Connected to MongoDB
All Sessions deleted
Load sessions from JSON file
Session saved successfully
Session saved successfully
Session saved successfully
Session saved successfully
Clean Speakers table
All Speakers deleted
Load Speakers from JSON file
Speaker saved successfully
Speaker saved successfully
Speaker saved successfully
Speaker saved successfully
```

8. Next, view your Cosmos DB instance in the Azure portal and see that it now contains two collections.

Add Collection Refresh Move Delete Account Data Explorer

Status Online	Read Locations East US, West US
Resource group (change) fabmedical-sol	Write Location East US
	URI https://fabmedical-sol2.documents.azure.com:443

Collections

ID	DATABASE	THROUGHPUT (RU/S)
sessions	contentdb	1000
speakers	contentdb	1000

Regions

Region Configuration FABMEDICAL-SOL2

Task 6: Test the application in a browser

In this task, you will verify that you can browse to the web service you have deployed and view the speaker and content information exposed by the API service.

1. From the Kubernetes management dashboard, in the navigation menu, select the **Services** view under **Discovery and Load Balancing**.
2. In the list of services, locate the external endpoint for the web service and select this hyperlink to launch the application.

Services						
Name	Labels	Cluster IP	Internal endpoints	External endpoints	Age	⋮
web	app: web	10.0.179.86	web:80 TCP web:32728 TCP	52.174.124.171:80 ↗	an hour	⋮
api	app: api	10.0.234.90	api:3001 TCP api:0 TCP	-	2 hours	⋮
kubernetes	component: apis... provider: kubern...	10.0.0.1	kubernetes:443 TCP kubernetes:0 TCP	-	7 hours	⋮

3. You will see the web application in your browser and be able to select the Speakers and Sessions links to view those pages without errors. The lack of errors means that the web application is correctly calling the API service to show the details on each of those pages.





SEPTEMBER 14-17, 2017

Monterey Conference Center
Monterey, California

sessions



Task 7: Configure Continuous Delivery to the Kubernetes Cluster

In this task, you will use GitHub Actions workflows to automate the process for deploying the web image to the AKS cluster. You will update the DevOps Pipeline and configure a deployment stage so that when new images are pushed to the ACR, the pipeline deploys the image to the AKS cluster.

1. Navigate to the `.github/workflows` folder of the git repository, and open the `content-web.yml` workflow using `vi`:

```
cd ~/MCW-Cloud-native-applications/Hands-on\ lab/lab-files/developer/.github/workflows
vi content-web.yml
```

2. You will add a second job to the bottom of the `content-web.yml` workflow. Paste the following at the end of the file:

Note: Be careful to check your indenting when pasting. The `build-and-push-helm-chart` node should be indented with 2 spaces and line up with the node for the `build-and-publish-docker-image` job.

```
build-and-push-helm-chart:
  name: Build and Push Helm Chart
  runs-on: ubuntu-latest
  needs: [build-and-publish-docker-image]
  steps:
    # Checkout the repo
    - uses: actions/checkout@master

    - name: Helm Install
      uses: azure/setup-helm@v1

    - name: Helm Repo Add
      run: |
        helm repo add ${{ env.containerRegistryName }} https://${{ env.containerRegistry }}/helm/v1/repo --username ${{ secrets.ACR_USERNAME }} --password ${{ secrets.ACR_PASSWORD }} --ca-insecure
        HELM_EXPERIMENTAL_OCI: 1

    - name: Helm Chart Save
      run: |
        cd ./content-web/charts/web

        helm chart save . content-web:v${{ env.tag }}
        helm chart save . ${{ env.containerRegistry }}/helm/content-web:v${{ env.tag }}

    # List out saved charts
    helm chart list
env:
```

```

HELM_EXPERIMENTAL_OCI: 1

- name: Helm Chart Push
  run: |
    helm registry login ${{ env.containerRegistry }} --username ${{ secrets.ACR_USERNAME }} --password ${{ secrets.ACR_PASSWORD }}
    helm chart push ${{ env.containerRegistry }}/helm/content-web:v${{ env.tag }}
  env:
    HELM_EXPERIMENTAL_OCI: 1</code></pre></li>

```

3. Save the file.

4. In the Azure Cloud Shell, use the following command to output the `~/.kube/config` file that contains the credentials for the AKS cluster.

```
cat ~/.kube/config
```

5. In GitHub, return to the `Fabmedical` repository screen, select the **Settings** tab, select **Secrets** from the left menu, and then click the **Add secret** button.

Secrets / New secret

Name

Value

NWWd:

QgUINBIFBSSVZBVEUgS0VZLS0tLS0K

token:

f6

9

Add secret

7. Now return to edit the `content-web.yml` workflow and paste the following at the end of the file.

Note: Be careful to check your indenting when pasting. The `aks-deployment` node should be indented with 2 spaces and the `uses` step should be indented with 4 spaces.

```

aks-deployment:
  name: AKS Deployment
  runs-on: ubuntu-latest
  needs: [build-and-publish-docker-image,build-and-push-helm-chart]
  steps:
    # Checkout the repo
    - uses: actions/checkout@master

```

```

- name: Helm Install
  uses: azure/setup-helm@v1

- name: kubeconfig
  run: echo "${{ secrets.KUBECONFIG }}" &gt;&gt; kubeconfig

- name: Helm Repo Add
  run: |
    helm repo add ${{ env.containerRegistry }} https://${{ env.containerRegistry }}/helm/v1/repo --username ${{ secrets.ACR_USERNAME }} --password ${{ secrets.ACR_PASSWORD }}
    helm repo update

env:
  HELM_EXPERIMENTAL_OCI: 1

- name: Helm Upgrade
  run: |
    helm registry login ${{ env.containerRegistry }} --username ${{ secrets.ACR_USERNAME }} --password ${{ secrets.ACR_PASSWORD }}
    helm chart pull ${{ env.containerRegistry }}/helm/content-web:v${{ env.tag }}
    helm chart export ${{ env.containerRegistry }}/helm/content-web:v${{ env.tag }} --destination ./upgrade
    helm upgrade web ./upgrade/web

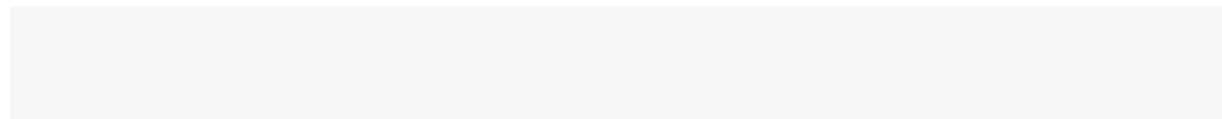
env:
  KUBECONFIG: './kubeconfig'
  HELM_EXPERIMENTAL_OCI: 1</code></pre></li>

```

8. Save the file.

9. On the `content-web` workflow, select **Run workflow** and manually trigger the workflow to execute.

10. Selecting the currently running workflow will display its status.



Task 8: Review Azure Monitor for Containers

In this task, you will access and review the various logs and dashboards made available by Azure Monitor for Containers.

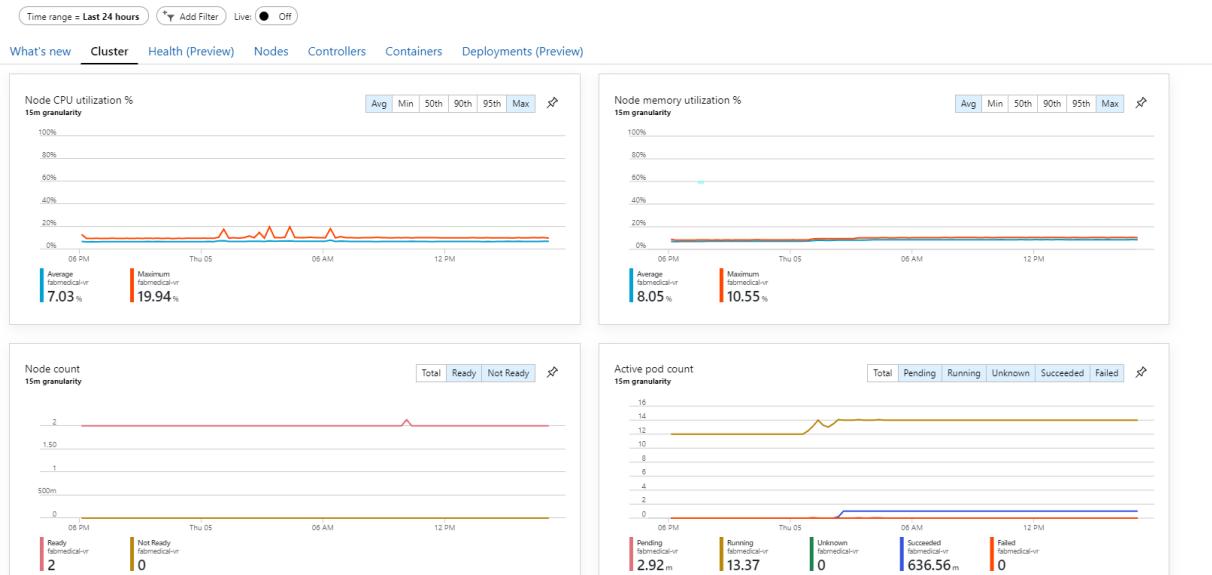
- From the Azure Portal, select the resource group you created named `fabmedical-SUFFIX`, and then select your AKS cluster.

NAME	TYPE	LOCATION	...
fabmedicald-ns	Network security group	East US	...
fabmedical	Container registry	East US	...
fabmedical-v	Virtual machine	East US	...
fabmedical-k	Kubernetes service	East US	...

- From the Monitoring blade, select **Insights**.



- Review the various available dashboards and a deeper look at the various metrics and logs available on the Cluster, Cluster Nodes, Cluster Controllers, and deployed Containers.



- To review the Containers dashboards and see more detailed information about each container, select the containers tab.

Time range = Last 24 hours

What's new Cluster Health (Preview) Nodes Controllers **Containers** Deployments (Preview)

Search by name... Metric: CPU Usage (millicores) ▾ Min Avg 50th 90th 95th Max

NAME	STATUS	95TH %	POD	NODE	RESTARTS	UPTIME	TREND 95TH % (1 BAR = 1H)
omsagent	Ok	9%	omsagent-rs-67569...	aks-agentpool-216...	0	23 hours	
omsagent	Ok	6%	omsagent-fq7jm...	aks-agentpool-216...	0	23 hours	
omsagent	Ok	5%	omsagent-c25w9...	aks-agentpool-216...	0	22 hours	
tunnel-front	Ok	5%	tunnelfront-867765...	aks-agentpool-216...	0	23 hours	
heapster-nanny	Ok	1%	heapster-5fc7766c9...	aks-agentpool-216...	0	23 hours	
heapster	Ok	1%	heapster-5fc7766c9...	aks-agentpool-216...	0	23 hours	
main	Ok	0.5%	main-5f9333333...	kubernetes-dashbo...	0	23 hours	
kubedns	Ok	0.3%	kube-dns-v20-864b...	aks-agentpool-216...	0	23 hours	
kubedns	Ok	0.3%	kube-dns-v20-864b...	aks-agentpool-216...	0	23 hours	
kube-proxy	Ok	0.2%	kube-proxy-mcbk...	aks-agentpool-216...	0	23 hours	
kube-proxy	Ok	0.2%	kube-proxy-cz6vc...	aks-agentpool-216...	0	23 hours	
healthz	Ok	0.2%	healthz-5f8d453...	aks-agentpool-216...	0	23 hours	
healthz	Ok	0.2%	healthz-5f8d453...	aks-agentpool-216...	0	23 hours	
sidecar	Ok	0.1%	sidecar-5f8d453...	aks-agentpool-216...	0	23 hours	
sidecar	Ok	0.1%	sidecar-5f8d453...	aks-agentpool-216...	0	23 hours	

27 items

» omsagent Container
View live data (preview)
View in analytics
Container Name: omsagent
Container ID: a1e7ea291c1840641c6a13464ccf05cfbecd6dbf5876a17d41c0aa80eef229970a
Container Status: running
Container Status Reason: -
Image: azuremonitor/containerinsights/ciprod
Image Tag: ciprod11012019
Container Creation Time Stamp: 12/4/2019, 6:03:57 PM
Start Time: 12/4/2019, 6:03:57 PM
Finish Time: -

5. Now filter by container name and search for the web containers, you will see all the containers created in the Kubernetes cluster with the pod names. You can compare the names with those in the kubernetes dashboard.

Time range = Last 24 hours

What's new Cluster Health (Preview) Nodes Controllers **Containers** Deployments (Preview)

Search by name... Metric: CPU Usage (millicores) ▾ Min Avg 50th 90th 95th Max

NAME	STATUS	95TH %	POD	NODE	RESTARTS	UPTIME	TREND 95TH % (1 BAR = 1H)
web	Ok	0%	web-6dbc6d4d56...	aks-agentpool-216...	0	12 hours	
web	Done	-	web-c7cf08c9d-71kg...	aks-agentpool-216...	0	?	
web	Unk	-	web-696854dd8d-fl...	aks-agentpool-216...	0	?	
web	Unk	-	web-5b6dcfcff-p27c...	aks-agentpool-216...	0	?	
web	Unk	-	web-5bbfd59965-m...	aks-agentpool-216...	0	?	

5 of 27 items matching

» web Container
View live data (preview)
View in analytics
Container Name: web
Container ID: e4033082d73d26f1da6b341d5f1307ff1b09654ef78a973fc48f11f3fa4320d3
Container Status: running

6. By default, the CPU Usage metric will be selected displaying all cpu information for the selected container, to switch to another metric open the metric dropdown list and select a different metric.

Time range = Last 24 hours

What's new Cluster Health (Preview) Nodes Controllers **Containers** Deployments (Preview)

Search by name... Metric: CPU Usage (millicores) ▾ CPU Usage (millicores) Memory working set Memory Rss

NAME	STATUS	95TH %	POD	NODE	RESTARTS
web	Ok	0%	web-6dbc6d4d56...	aks-agentpool-216...	0

Min Avg 50th 90th 95th Max

7. Upon selecting any pod, all the information related to the selected metric will be displayed on the right panel, and that would be the case when selecting any other metric, the details will be displayed on the right panel for the selected pod.

»



web

Container

[View live data \(preview\)](#)

[View in analytics](#)

▼

Container Name

web

Container ID

e4033082d73d26f1da6b341d5f1307ff1b09654ef78a973fc48f
11f3fa4320d3

Container Status

running

Container Status Reason

-

Image

content-web

Image Tag

latest

Container Creation Time Stamp

12/5/2019, 4:18:30 AM

Start Time

12/5/2019, 4:18:30 AM

Finish Time

-

CPU Limit

1900 mc

CPU Request

1000 mc

Memory Limit

4.45 GB

Memory Request

128 MB

Last reported

16 secs ago

▷ Environment Variables

8. To display the logs for any container simply select it and view the right panel and you will find **View container logs** option which will list all logs for this specific container.

```

let startDateTime = datetime('2019-12-04T17:00:00.000Z');
let endDateTime = datetime('2019-12-05T17:08:46.332Z');

let ContainerName = 'Kubernetes';
let TimeGenerated >= startDateTime and TimeGenerated < endDateTime
where ContainerName == '392cfe8a-1716-11ea-b01c-1a15b0b45cd5/web'
where ClusterId == '/subscriptions/0a23088d-13bf-4331-9e70-23fcc08720ef/resourceGroups/fabmedical-vn/providers/Microsoft.ContainerService/managedClusters/fabmedical-vn'
ContainerLog
| where TimeGenerated > startDateTime and TimeGenerated < endDateTime
| where ContainerID in (ContainerIDList)
| project LogEntrySource, LogEntry, TimeGenerated, Computer, Image, Name, ContainerID
    
```

TimeGenerated [UTC]	LogEntrySource	LogEntry	Computer	Image	Name	ContainerID
12/5/2019, 4:18:30.652 AM	stdout	Running	aks-agentpool-21681325-0		e4033082d73d26ffda6b341d5f1307fffb09654ef78a973fc48ff1f3fa4320d3	

9. For each log entry you can display more information by expanding the log entry to view the below details.

TimeGenerated [UTC]	LogEntrySource	LogEntry	Computer	Image	Name	ContainerID
12/5/2019, 4:18:30.652 AM	stdout	Running	aks-agentpool-21681325-0		e4033082d73d26ffda6b341d5f1307fffb09654ef78a973fc48ff1f3fa4320d3	
	LogEntrySource	stdout				
	LogEntry	Running				
	TimeGenerated [UTC]	2019-12-05T04:18:30.652Z				
	Computer	aks-agentpool-21681325-0				
	ContainerID	e4033082d73d26ffda6b341d5f1307fffb09654ef78a973fc48ff1f3fa4320d3				

Exercise 3: Scale the application and test HA

Duration: 20 minutes

At this point, you have deployed a single instance of the web and API service containers. In this exercise, you will increase the number of container instances for the web service and scale the front-end on the existing cluster.

Task 1: Increase service instances from the Kubernetes dashboard

In this task, you will increase the number of instances for the API deployment in the Kubernetes management dashboard. While it is deploying, you will observe the changing status.

- From the navigation menu, select **Workloads -> Deployments**, and then select the **API** deployment.

- Select **SCALE**.

☰ Workloads > Deployments > api

☰ SCALE EDIT DELETE

3. Change the number of pods to 2, and then select **OK**.

Scale a resource

deployment api will be updated to reflect the desired replicas count.

Desired replicas *	Actual replicas
2	1

i This action is equivalent to: `kubectl scale -n default deployment api --replicas=2`

Scale **Cancel**

Note: If the deployment completes quickly, you may not see the deployment Waiting states in the dashboard, as described in the following steps.

4. From the **Replica Set** view for the **API**, you will see it is now deploying and that there is one healthy instance and one pending instance.

☰ Workloads > Replica Sets > api-2547917202

☰ LOGS ⚡ SCALE ⚡ EDIT ⚡ DELETE

Cluster	Name: api-2547917202
Namespaces	Namespace: default
Nodes	Labels: app: api pod-template-hash: 2547917202
Persistent Volumes	Annotations: deployment.kubernetes.io/desired-replicas: 2 deployment.kubernetes.io/max-replicas: 3 deployment.kubernetes.io/revision: 2
Roles	Creation time: 2018-01-25T22:41
Storage Classes	Selector: app: api pod-template-hash: 2547917202
Namespace	Images: fabmedical .azuredcr.io/fabmedical/content-api
default	Status
Overview	Pods: 2 created, 2 desired Pods status: 1 pending, 1 running
Workloads	
Daemon Sets	
Deployments	
Jobs	
Pods	api-2547917202-51g... k8s-agent-b882c7f1-1 Waiting: Contai... 0 3 seconds ⚡ ⚡
Replica Sets	api-2547917202-glw... k8s-agent-b882c7f1-0 Running 0 13 hours ⚡ ⚡
Replication Controllers	
Stateful Sets	

5. From the navigation menu, select **Deployments** from the list. Note that the api service has a pending status indicated by the grey timer icon, and it shows a pod count 1 of 2 instances (shown as 1/2).

The screenshot shows the Kubernetes Deployment page. On the left, there's a sidebar with a navigation menu under 'Workloads'. The 'Deployments' option is selected. The main panel is titled 'Deployments' and lists two entries:

Name	Labels	Pods	Age	Images
web	app: web	1 / 1	12 hours	fabmedical ...
api	app: api	1 / 2	13 hours	fabmedical ...

The 'api' row is highlighted with a red outline. The 'api' service has one healthy pod and one pod in progress.

6. From the Navigation menu, select **Workloads**. From this view, note that the health overview in the right panel of this view. You will see the following:

- One deployment and one replica set are each healthy for the api service.
- One replica set is healthy for the web service.
- Three pods are healthy.

7. Navigate to the web application from the browser again. The application should still work without errors as you navigate to Speakers and Sessions pages

- Navigate to the `/stats` page. You will see information about the environment including:
 - **webTaskId**: The task identifier for the web service instance.
 - **taskId**: The task identifier for the API service instance.
 - **hostName**: The hostname identifier for the API service instance.
 - **pid**: The process id for the API service instance.
 - **mem**: Some memory indicators returned from the API service instance.
 - **counters**: Counters for the service itself, as returned by the API service instance.
 - **uptime**: The up time for the API service.
- Refresh the page in the browser, and you can see the **hostName** change between the two API service instances. The letters after `api-{number}-` in the hostname will change.

Task 2: Increase service instances beyond available resources

In this task, you will try to increase the number of instances for the API service container beyond available resources in the cluster. You will observe how Kubernetes handles this condition and correct the problem.

1. From the navigation menu, select **Deployments**. From this view, select the **API** deployment.

2. Configure the deployment to use a fixed host port for initial testing. Select **Edit**.

```
1  {
2    "kind": "Deployment",
3    "apiVersion": "extensions/v1beta1",
4    "metadata": {
5      "name": "api",
6      "namespace": "default",
7      "selfLink": "/apis/extensions/v1beta1/namespaces/default/deployments/api",
8      "uid": "9db5f84a-31ca-11e8-aa11-0a58ac1f01",
9      "resourceVersion": "64695",
10     "generation": 1,
11     "creationTimestamp": "2018-03-27T14:24:53Z",
12     "labels": {
13       "k8s-app": "api"
14     },
15     "annotations": {
16       "deployment.kubernetes.io/revision": "1"
17     }
}
```

CANCEL COPY UPDATE

4. Paste the contents into the text editor of your choice (*notepad is shown here, macOS users can useTextEdit*).

```

Untitled - Notepad
File Edit Format View Help

{
  "kind": "Deployment",
  "apiVersion": "extensions/v1beta1",
  "metadata": {
    "name": "api",
    "namespace": "default",
    "selfLink": "/apis/extensions/v1beta1/namespaces/default/deployments/api",
    "uid": "9db5f84a-31ca-11e8-aa11-0a58ac1f0587",
    "resourceVersion": "64695",
    "generation": 1,
    "creationTimestamp": "2018-03-27T14:24:53Z",
    "labels": {
      "k8s-app": "api"
    },
    "annotations": {
      "deployment.kubernetes.io/revision": "1"
    }
  },
  "spec": {
    "replicas": 1,
    "selector": {
      "matchLabels": {
        "k8s-app": "api"
      }
    }
  }
}

```

5. Scroll down about halfway to find the node `$.spec.template.spec.containers[0]`, as shown in the screenshot below.

[!Screenshot of the deployment JSON code, with the `\$.spec.template.spec.containers\[0\]` section highlighted.](#)

6. The containers spec has a single entry for the API container at the moment. You will see that the name of the container is `api` - this is how you know you are looking at the correct container spec.

- Add the following JSON snippet below the `name` property in the container spec:

```

"ports": [
  {
    "containerPort": 3001,
    "hostPort": 3001
  }
],

```

- Your container spec should now look like this:

[!Screenshot of the deployment JSON code, with the `\$.spec.template.spec.containers\[0\]` section highlighted, showing the updated values for `containerPort` and `hostPort`, both set to port 3001.](#)

1. Copy the updated JSON document from notepad into the clipboard. Return to the Kubernetes dashboard, which should still be viewing the `api` deployment.

- Select **Edit**.



- Paste the updated JSON document.
- Select **Update**.



2. From the API deployment view, select **Scale**.



3. Change the number of pods to **4** and select **OK**.

Scale a Deployment

Resource api will be updated to reflect the desired count.
Current status: 2 created, 4 desired.

Desired number of pods

4 X

CANCEL OK

- From the navigation menu, select **Services** view under **Discovery and Load Balancing**. Select the **api** service from the **Services** list. From the api service view, you will see it has two healthy instances and two unhealthy (or possibly pending depending on timing) instances.

The screenshot shows the Kubernetes Services view for the 'api' service. The left sidebar includes options like Storage Classes, Namespace (default), Overview, Workloads, Cron Jobs, Daemon Sets, Deployments, Jobs, Pods, Replica Sets, Replication Controllers, Stateful Sets, Ingresses, Services (selected), Config and Storage, Config Maps, Persistent Volume Claims, Secrets, Settings, and About. The main content area displays the 'api' service details, endpoints, and pods.

Details:

Name	Namespace	Labels	Creation Time	Type	Session Affinity	Connection	Cluster IP	Internal endpoints
api	default	k8s-app: api	2020-07-13T19:11 UTC	ClusterIP	None		10.0.108.133	api:3001 TCP

Endpoints:

Host	Ports (Name, Port, Protocol)	Node	Ready
10.244.0.12	tcp-3001-3001-k75q2, 3001, TCP	aks-agentpool-23575493-vmss000000	true
10.244.1.9	tcp-3001-3001-k75q2, 3001, TCP	aks-agentpool-23575493-vmss000001	true

Pods:

Name	Node	Status	Restarts	Age	⋮
api-64956d8b56-b59km		Pending	0	6 minutes	⋮
api-64956d8b56-fdkwr	aks-agentpool-23575493-vmss000000	Running	0	7 minutes	⋮
api-64956d8b56-hwdk	aks-agentpool-23575493-vmss000001	Running	0	7 minutes	⋮
api-64956d8b56-itvcm		Pending	0	6 minutes	⋮

- After a few minutes, select **Workloads** from the navigation menu. From this view, you should see an alert reported for the api deployment.

Workloads

default

Overview

Workloads

- Daemon Sets
- Deployments
- Jobs
- Pods
- Replica Sets
- Replication Controllers
- Stateful Sets

Discovery and Load Balancing

- Ingresses
- Services

Config and Storage

- Config Maps
- Persistent Volume Claims
- Secrets

Deployments

Name	Labels	Pods	Age	Images
web	app: web	1 / 1	12 hours	fabmedical...
api	app: api	2 / 4	14 hours	fabmedical...

No nodes are available that match all of the following predicates:: PodFitsHostPorts (2), PodToleratesNodeTaints (1).

Pods

Name	Node	Status	Restarts	Age
api-2547917202-5xcjk		Pending	0	57 seconds
api-2547917202-kjrp7		Pending	0	57 seconds
api-2547917202-51g...	k8s-agent-b882c7f1-1	Running	0	17 minutes
web-1272019779-f9...	k8s-agent-b882c7f1-0	Running	0	11 hours

Note: This message indicates that there were not enough available resources to match the requirements for a new pod instance. In this case, this is because the instance requires port 3001, and since there are only 2 nodes available in the cluster, only two api instances can be scheduled. The third and fourth pod instances will wait for a new node to be available that can run another instance using that port.

6. Reduce the number of requested pods to 2 using the **Scale** button.

7. Almost immediately, the warning message from the Workloads dashboard should disappear, and the API deployment will show "2/2" pods are running.

Workloads

default

Overview

Workloads

- Daemon Sets
- Deployments
- Jobs
- Pods
- Replica Sets
- Replication Controllers
- Stateful Sets

Discovery and Load Balancing

- Ingresses
- Services

Config and Storage

- Config Maps
- Persistent Volume Claims
- Secrets

Deployments

Name	Labels	Pods	Age	Images
web	app: web	1 / 1	12 hours	fabmedical...
api	app: api	2 / 2	14 hours	fabmedical...

Pods

Name	Node	Status	Restarts	Age
api-2547917202-51g...	k8s-agent-b882c7f1-1	Running	0	22 minutes
web-1272019779-f9...	k8s-agent-b882c7f1-0	Running	0	11 hours
api-2547917202-glw...	k8s-agent-b882c7f1-0	Running	0	13 hours

Replica Sets

Name	Labels	Pods	Age	Images
web	app: web	1 / 1	12 hours	fabmedical...
api	app: api	2 / 2	14 hours	fabmedical...

Task 3: Restart containers and test HA

In this task, you will restart containers and validate that the restart does not impact the running service.

- From the navigation menu on the left, select **Services** view under **Discovery and Load Balancing**. From the **Services** list, select the external endpoint hyperlink for the web service, and visit the stats page by adding / stats to the URL. Keep this open and handy to be refreshed as you complete the steps that follow.

Name	Labels	Cluster IP	Internal endpoints	External endpoints	Age	⋮
web	app: web	10.0.179.86	web:80 TCP web:32728 TCP	52.174.124.171:80 ↗	an hour	⋮
api	app: api	10.0.234.90	api:3001 TCP api:0 TCP	-	2 hours	⋮
kubernetes	component: apis... provider: kubern...	10.0.0.1	kubernetes:443 TCP kubernetes:0 TCP	-	7 hours	⋮

CONTOSO NEURO 2017

Speakers Sessions



Stats

webTaskId	18
taskId	18
hostName	api-2547917202-glwz7
pid	18
mem	{"rss":36429824,"heapTotal":19582720,"heapUsed":13460680,"external":232284}
counters	{"stats":3,"speakers":2,"sessions":1}
uptime	49422.159

- From the navigation menu, select **Workloads -> Deployments**. From Deployments list, select the **API** deployment.

The screenshot shows the Kubernetes Deployment view. On the left, there's a navigation sidebar with sections like Cluster, Namespaces, Nodes, Persistent Volumes, Roles, Storage Classes, Namespace (set to default), Overview, Workloads (with Deployments selected, indicated by a red arrow), Jobs, Pods, Replica Sets, Replication Controllers, and Stateful Sets. The main area is titled 'Deployments' and lists two entries:

Name	Labels	Pods	Age	Images
web	app: web	1 / 1	13 hours	fabmedical ...
api	app: api	2 / 2	14 hours	fabmedical ...

3. From the API deployment view, select **Scale** and from the dialog presented, and enter `4` for the desired number of pods. Select **OK**.
4. From the navigation menu, select **Workloads -> Replica Sets**. Select the `api` replica set, and from the Replica Set view, you will see that two pods cannot deploy.

Details

Name: api-2547917202
Namespace: default
Labels: app: api, pod-template-hash: 2547917202
Annotations: deployment.kubernetes.io/desired-replicas: 4, deployment.kubernetes.io/max-replicas: 5, deployment.kubernetes.io/revision: 2
Creation time: 2018-01-25T22:41
Selector: app: api, pod-template-hash: 2547917202
Images: fabmedical.azurecr.io/fabmedical/content-api
Status
Pods: 4 created, 4 desired
Pods status: 2 pending, 2 running

Name	Node	Status	Restarts	Age	⋮
api-2547917202-4x1...		Pending	0	7 seconds	⋮
api-2547917202-51g...	k8s-agent-b882c7f1-1	Running	0	39 minutes	⋮
api-2547917202-glw...	k8s-agent-b882c7f1-0	Running	0	13 hours	⋮
api-2547917202-hjzsf		Pending	0	7 seconds	⋮

5. Return to the browser tab with the web application stats page loaded. Refresh the page over and over. You will not see any errors, but you will see the api host name change between the two api pod instances periodically. The task id and pid might also change between the two api pod instances.

webTaskId	18
taskId	18
hostName	api-2547917202-glwz7
pid	18
mem	{"rss":36429824,"heapTotal":19582720,"heapUsed":13544928,"external":232284}
counters	{"stats":6,"speakers":2,"sessions":1}
uptime	50570.163

webTaskId	18
taskId	18
hostName	api-2547917202-51gjc
pid	18
mem	{"rss":36376576,"heapTotal":19582720,"heapUsed":13479320,"external":237370}
counters	{"stats":5,"speakers":3,"sessions":1}
uptime	3012.69

6. After refreshing enough times to see that the hostName value is changing, and the service remains healthy, return to the Replica Sets view for the API. From the navigation menu, select **Replica Sets** under **Workloads** and select the API replica set.

7. From this view, take note that the hostName value shown in the web application stats page matches the pod names for the pods that are running.

Pods						
	Name	Node	Status	Restarts	Age	
!	api-2547917202-4x1...		Pending	0	7 seconds	
	No nodes are available that match all of the following predicates:: PodFitsHostPorts (2), PodToleratesNodeTaints (1).					
✓	api-2547917202-51g...	k8s-agent-b882c7f1-1	Running	0	39 minutes	
✓	api-2547917202-glw...	k8s-agent-b882c7f1-0	Running	0	13 hours	
!	api-2547917202-hjzsf		Pending	0	7 seconds	
	No nodes are available that match all of the following predicates:: PodFitsHostPorts (2), PodToleratesNodeTaints (1).					

8. Note the remaining pods are still pending, since there are not enough port resources available to launch another instance. Make some room by deleting a running instance. Select the context menu and choose **Delete** for one of the healthy pods.

Pods						
	Name	Node	Status	Restarts	Age	
!	api-2547917202-4x1...		Pending	0	30 minutes	
	No nodes are available that match all of the following predicates:: PodFitsHostPorts (2), PodToleratesNodeTaints (1).					
✓	api-2547917202-51g...	k8s-agent-b882c7f1-1	Running	0	an hour	
✓	api-2547917202-glw...	k8s-agent-b882c7f1-0	Running	0	14 hours	
!	api-2547917202-hjzsf		Pending	0	30 minutes	
	No nodes are available that match all of the following predicates:: PodFitsHostPorts (2), PodToleratesNodeTaints (1).					

9. Once the running instance is gone, Kubernetes will be able to launch one of the pending instances. However, because you set the desired size of the deploy to 4, Kubernetes will add a new pending instance. Removing a running instance allowed a pending instance to start, but in the end, the number of pending and running instances is unchanged.

Pods						
Name	Node	Status	Restarts	Age		
api-2547917202-4x1...	k8s-agent-b882c7f1-1	Running	0	37 minutes		
api-2547917202-glw...	k8s-agent-b882c7f1-0	Running	0	14 hours		
api-2547917202-hjzsf		Pending	0	37 minutes		
	No nodes are available that match all of the following predicates:: PodFitsHostPorts (2), PodToleratesNodeTaints (1).					
api-2547917202-nqq...		Pending	0	58 seconds		
	No nodes are available that match all of the following predicates:: PodFitsHostPorts (2), PodToleratesNodeTaints (1).					

10. From the navigation menu, select **Deployments** under **Workloads**. From the view's Deployments list, select the **API** deployment.

11. From the API Deployment view, select Scale and enter **1** as the desired number of pods. Select **OK**.

Scale a Deployment

Resource api will be updated to reflect the desired count.
Current status: 4 created, 1 desired.

Desired number of pods

1

X

CANCEL

OK

12. Return to the web site's stats page in the browser and refresh while this is scaling down. You will notice that only one API host name shows up, even though you may still see several running pods in the API replica set view. Even though several pods are running, Kubernetes will no longer send traffic to the pods it has selected to scale down. In a few moments, only one pod will show in the API replica set view.

Details

- Name: api-2547917202
- Namespace: default
- Labels: app: api, pod-template-hash: 2547917202
- Annotations: deployment.kubernetes.io/desired-replicas: 1, deployment.kubernetes.io/max-replicas: 2, deployment.kubernetes.io/revision: 2
- Creation time: 2018-01-25T22:41
- Selector: app: api, pod-template-hash: 2547917202
- Images: fabmedical.azurecr.io/fabmedical/content-api

Status

Pods: 1 running

Name	Node	Status	Restarts	Age
api-2547917202-glw...	k8s-agent-b882c7f1-0	Running	0	14 hours

Pods

13. From the navigation menu, select **Workloads**. From this view, note that there is only one API pod now.

Deployments

Name	Labels	Pods	Age	Images
web	app: web	1 / 1	14 hours	fabmedical...
api	app: api	1 / 1	15 hours	fabmedical...

Pods

Name	Node	Status	Restarts	Age
web-1272019779-f9...	k8s-agent-b882c7f1-0	Running	0	12 hours
api-2547917202-glw...	k8s-agent-b882c7f1-0	Running	0	14 hours

Replica Sets

Name	Labels	Pods	Age	Images
web-1272019779	app: web, pod-template...	1 / 1	12 hours	fabmedical...

Exercise 4: Working with services and routing application traffic

Duration: 45 minutes

In the previous exercise, we introduced a restriction to the scale properties of the service. In this exercise, you will configure the api deployments to create pods that use dynamic port mappings to eliminate the port resource constraint during scale activities.

Kubernetes services can discover the ports assigned to each pod, allowing you to run multiple instances of the pod on the same agent node --- something that is not possible when you configure a specific static port (such as 3001 for the API service).

Task 1: Scale a service without port constraints

In this task, we will reconfigure the API deployment so that it will produce pods that choose a dynamic hostPort for improved scalability.

1. From the navigation menu select **Deployments** under **Workloads**. From the view's Deployments list, select the **API** deployment.
2. Select **Edit**.
3. From the **Edit a Deployment** dialog, do the following:
 - Scroll to the first spec node that describes replicas as shown in the screenshot. Set the value for replicas to **4**.
 - Within the replicas spec, beneath the template node, find the `api` containers spec. Remove the hostPort entry for the API container's port mapping. The screenshot below shows the desired configuration after editing.



```
19: "spec": {
20:   "replicas": 4,
21:   "selector": [
22:     {
23:       "matchLabels": {
24:         "k8s-app": "api"
25:       }
26:     },
27:     {
28:       "template": {
29:         "metadata": {
30:           "name": "api",
31:           "creationTimestamp": null,
32:           "labels": {
33:             "k8s-app": "api"
34:           }
35:         },
36:         "spec": {
37:           "containers": [
38:             {
39:               "name": "api",
40:               "image": "fabmedicalsol.azurecr.io/fabmedical/content-ap
41:               "ports": [
42:                 {
43:                   "containerPort": 3001,
44:                   "protocol": "TCP"
45:                 }
46:               ],
47:               "resources": {
48:                 "limits": {
49:                   "cpu": "0.1",
50:                   "memory": "128Mi"
51:                 },
52:                 "requests": {
53:                   "cpu": "0.1",
54:                   "memory": "128Mi"
55:                 }
56:               }
57:             }
58:           ]
59:         }
60:       }
61:     }
62:   }
63: }
```

4. Select **Update**. New pods will now choose a dynamic port.

5. The API service can now scale to 4 pods since it is no longer constrained to an instance per node -- a previous limitation while using port 3001 .

The screenshot shows the AKS portal interface. At the top, there's a navigation bar with 'Workloads > Replica Sets > api-4178426672'. To the right of the navigation are buttons for 'LOGS', 'SCALE', 'EDIT', and 'DELETE'. On the left, a sidebar menu lists 'Cluster', 'Namespaces', 'Nodes', 'Persistent Volumes', 'Roles', 'Storage Classes', and a 'Namespace' dropdown set to 'default'. Below this is a 'Workloads' section with options: 'Daemon Sets', 'Deployments', 'Jobs', 'Pods', 'Replica Sets' (which is selected and highlighted in blue), and 'Replication Controllers'. Under 'Replica Sets', there's also a 'Stateful Sets' option. The main content area has a 'Labels' section with 'app: api' and 'pod-template-hash: 4178426672'. It also shows 'Annotations' like 'deployment.kubernetes.io/desired-replicas: 4' and 'deployment.kubernetes.io/max-replicas: 5'. Below that are 'Creation time: 2018-01-26T17:06', 'Selector: app: api', and 'Images: fabmedical.azurecr.io/fabmedical/content-api'. A 'Status' section indicates 'Pods: 4 running'. The bottom half of the screen is a table titled 'Pods' with columns: Name, Node, Status, Restarts, and Age. There are four rows, each representing a pod named 'api-4178426672-[long ID]' running on 'k8s-agent-b882c7f1-1' with 0 restarts and an age of 5 minutes.

Name	Node	Status	Restarts	Age	Actions
api-4178426672-45lxx	k8s-agent-b882c7f1-1	Running	0	5 minutes	⋮
api-4178426672-76l...	k8s-agent-b882c7f1-1	Running	0	5 minutes	⋮
api-4178426672-sph...	k8s-agent-b882c7f1-0	Running	0	5 minutes	⋮
api-4178426672-vs...	k8s-agent-b882c7f1-1	Running	0	5 minutes	⋮

6. Return to the browser and refresh the stats page. You should see all 4 pods serve responses as you refresh.

Task 2: Update an external service to support dynamic discovery with a load balancer

In this task, you will update the web service so that it supports dynamic discovery through the Azure load balancer.

- From the navigation menu, select **Deployments** under **Workloads**. From the view's Deployments list, select the **web** deployment.
- Select **Edit**.
- From the **Edit a Deployment** dialog, scroll to the web containers spec as shown in the screenshot. Remove the `hostPort` entry for the web container's port mapping.

Edit a Deployment

```
 24      app: "web"
 25    },
 26  },
 27  "template": {
 28    "metadata": {
 29      "name": "web",
 30      "creationTimestamp": null,
 31      "labels": {
 32        "app": "web"
 33      }
 34    },
 35    "spec": {
 36      "containers": [
 37        {
 38          "name": "web",
 39          "image": "fabmedicalisoli.azurecr.io/fabmedical/content-web:latest",
 40          "ports": [
 41            {
 42              "containerPort": 3000,
 43              "protocol": "TCP"
 44            }
 45          ],
 46          "env": [
 47            {
 48              "name": "CONTENT_API_URL",
 49              "value": "http://api:3001"
 50            }
 51          ]
 52        }
 53      ]
 54    }
 55  }
 56}
```

The code editor shows a deployment manifest. A red box highlights the container's ports configuration, specifically the port mapping from the container's port 3000 to the host's port 3000. Below the code editor are three buttons: CANCEL, COPY, and UPDATE.

4. Select **Update**.

5. From the web Deployments view, select **Scale**. From the dialog presented enter 4 as the desired number of pods and select **OK**.

6. Check the status of the scale out by refreshing the web deployment's view. From the navigation menu, select Deployments from under Workloads. Select the web deployment. From this view, you should see an error like that shown in the following screenshot.

Events

Message	Source
! 0/2 nodes are available: 2 Insufficient cpu.	default-scheduler
! 0/2 nodes are available: 2 Insufficient cpu.	default-scheduler

Like the API deployment, the web deployment used a fixed `hostPort`, and your ability to scale was limited by the number of available agent nodes. However, after resolving this issue for the web service by removing the `hostPort` setting, the web deployment is still unable to scale past two pods due to CPU constraints. The deployment is requesting more CPU than the web application needs, so you will fix this constraint in the next task.

Task 3: Adjust CPU constraints to improve scale

In this task, you will modify the CPU requirements for the web service so that it can scale out to more instances.

1. From the navigation menu, select **Deployments** under **Workloads**. From the view's Deployments list, select the web deployment.
2. Select **Edit**.
3. From the Edit a Deployment dialog, find the `cpu` resource requirements for the web container. Change this value to `125m`.



4. Select **Update** to save the changes and update the deployment.
5. From the navigation menu, select **Replica Sets** under **Workloads**. From the view's **Replica Sets** list select the web replica set.

6. When the deployment update completes, four web pods should be shown in running state.

Pods						
	Name	Node	Status	Restarts	Age	
✓	web-120118169-0nlfh	k8s-agent-b882c7f1-1	Running	0	a minute	
✓	web-120118169-86lrj	k8s-agent-b882c7f1-1	Running	0	a minute	
✓	web-120118169-fzztr	k8s-agent-b882c7f1-1	Running	0	a minute	
✓	web-120118169-rf06c	k8s-agent-b882c7f1-0	Running	0	a minute	

7. Return to the browser tab with the web application loaded. Refresh the stats page at / stats to watch the display update to reflect the different api pods by observing the host name refresh.

Task 4: Perform a rolling update

In this task, you will edit the web application source code to add Application Insights and update the Docker image used by the deployment. Then you will perform a rolling update to demonstrate how to deploy a code change.

1. Execute this command in Azure Cloud Shell to retrieve the instrumentation key for the `content-web` Application Insights resource:

```
az resource show -g fabmedical-[SUFFIX] -n content-web --resource-type "Microsoft.Insights/components" --query properties
```

Copy this value. You will use it later.

2. Update your starter files by pulling the latest changes from the Git repository:

```
cd ~/MCW-Cloud-native-applications/Hands-on\ lab/lab-files/infrastructure/content-web  
git pull
```

3. Install support for Application Insights.

```
npm install applicationinsights --save
```

4. Open the `app.js` file:

```
code app.js
```

5. Add the following lines immediately after `express` is instantiated:

```
const appInsights = require("applicationinsights");  
appInsights.setup("[YOUR APPINSIGHTS KEY]");  
appInsights.start();
```

Azure Cloud Shell

Bash

```
1 const express = require('express');
2 const http = require('http');
3 const path = require('path');
4
5 const app = express();
6
7 const appInsights = require("applicationinsights");
8 appInsights.setup("f6d88c82-
9 appInsights.start();
10
11 app.use(express.static(path.join(__dirname, 'dist/content-web')));
12 app.get('/config/content', function (req, res) {
13   const contentAniUrl = process.env.CONTENT_ANI_URL || "https://localhost:3001/".
```

6. Save changes and close the editor.

7. Push these changes to your repository so that GitHub Actions CI will build and deploy a new image.

```
git add .
git commit -m "Added Application Insights"
git push
```

8. Visit the `content-web` workflow for your GitHub repository and see the new image being deployed into your Kubernetes cluster.

9. While this update runs, return the Kubernetes management dashboard in the browser.

10. From the navigation menu, select **Replica Sets** under **Workloads**. From this view, you will see a new replica set for the web, which may still be in the process of deploying (as shown below) or already fully deployed.

Name	Labels	Pods	Age	Images
web-250784948	app: web pod-template...	0 / 2	0 seconds	fabmedical
web-3594286523	app: web pod-template...	4 / 3	a minute	fabmedical

11. While the deployment is in progress, you can navigate to the web application and visit the stats page at `/stats`. Refresh the page as the rolling update executes. Observe that the service is running normally, and tasks continue to be load balanced.



webTaskId	web-250784948-g9fk4
taskId	18
hostName	api-4178426672-f0dp8
pid	18
mem	{"rss":37785600,"heapTotal":19582720,"heapUsed":13256432,"external":211680}
counters	{"stats":4,"speakers":0,"sessions":0}
uptime	2559.844

Task 5: Configure Kubernetes Ingress

In this task you will setup a Kubernetes Ingress to take advantage of path-based routing and TLS termination.

1. Update your helm package list.

```
helm repo update
```

Note: If you get a "no repositories found." error, then run the following command. This will add back the official Helm "stable" repository.

```
helm repo add stable https://kubernetes-charts.storage.googleapis.com/
```

2. Install the ingress controller resource to handle ingress requests as they come in. The ingress controller will receive a public IP of its own on the Azure Load Balancer and be able to handle requests for multiple services over port 80 and 443.

```
helm install stable/nginx-ingress --namespace kube-system --set controller.replicaCount=2 --generate-name
```

3. From the Kubernetes dashboard, under **Discovery and Load Balancing**, select **Services**, then copy the IP Address for the **External endpoints** for the ingress-controller-nginx service.

Services					
Name	Labels	Cluster IP	Internal endpoints	External endpoints	Age
ingress-controller-nginx...	app: nginx-ingress chart: nginx-ingress-1.2. component: controller heritage: Helm release: ingress-control...	10.0.26.40	ingress-controller-nginx... ingress-controller-nginx... ingress-controller-nginx... ingress-controller-nginx...	51.144.226.102:80 51.144.226.102:443	a minute

Note: Alternately, you can find the IP using the following command in Azure Cloud Shell.

```
kubectl get svc --namespace kube-system
```

@Azure:~\$ kubectl get svc --namespace kube-system						
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)		AGE
healthmodel-replicaset-service	ClusterIP	10.0.154.232	<none>	25227/TCP		12h
ingress-controller-nginx-ingress-controller	LoadBalancer	10.0.26.40	51.144.226.102	80:31716/TCP,443:30115/TCP		8m
ingress-controller-nginx-ingress-default-backend	ClusterIP	10.0.231.84	<none>	80/TCP		8m
kube-dns	ClusterIP	10.0.0.10	<none>	53/UDP,53/TCP		12h
kubernetes-dashboard	ClusterIP	10.0.198.13	<none>	80/TCP		12h
metrics-server	ClusterIP	10.0.196.10	<none>	443/TCP		12h

4. Create a script to update the public DNS name for the IP.

```
code update-ip.sh
```

Paste the following as the contents and update the IP and SUFFIX values:

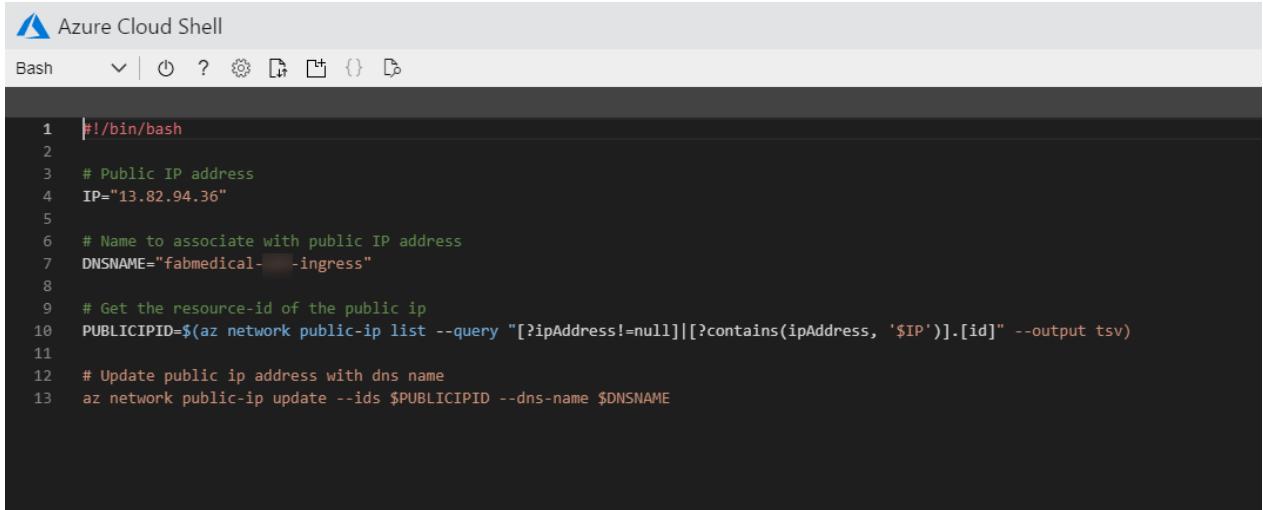
```
#!/bin/bash

# Public IP address
IP="[INGRESS PUBLIC IP]"

# Name to associate with public IP address
DNSNAME="fabmedical-[SUFFIX]-ingress"

# Get the resource-id of the public ip
PUBLICIPID=$(az network public-ip list --query "[?ipAddress!=null][?contains(ipAddress, '$IP')].[id]" --output tsv)

# Update public ip address with dns name
az network public-ip update --ids $PUBLICIPID --dns-name $DNSNAME
```



The screenshot shows the Azure Cloud Shell interface with a dark theme. At the top, there's a toolbar with icons for Bash, Copy, Help, and others. Below the toolbar, the command line shows the script content. The script is a bash script with variables for IP and DNSNAME, and it uses az CLI commands to find the public IP ID and update its DNS name.

```
1 #!/bin/bash
2
3 # Public IP address
4 IP="13.82.94.36"
5
6 # Name to associate with public IP address
7 DNSNAME="fabmedical-[SUFFIX]-ingress"
8
9 # Get the resource-id of the public ip
10 PUBLICIPID=$(az network public-ip list --query "[?ipAddress!=null][?contains(ipAddress, '$IP')].[id]" --output tsv)
11
12 # Update public ip address with dns name
13 az network public-ip update --ids $PUBLICIPID --dns-name $DNSNAME
```

Be sure to replace the following placeholders in the script:

- [INGRESS PUBLIC IP] : replace this with the IP Address copied previously.
- [SUFFIX] : replace this with the same SUFFIX value used previously for this lab

1. Save changes and close the editor.

2. Run the update script.

```
bash ./update-ip.sh
```

7. Verify the IP update by visiting the URL in your browser.

Note: It is normal to receive a 404 message at this time.

```
http://fabmedical-[SUFFIX]-ingress.[AZURE-REGION].cloudapp.azure.com/
```



8. Use helm to install `cert-manager`, a tool that can provision SSL certificates automatically from letsencrypt.org.

```
kubectl create namespace cert-manager  
  
kubectl label namespace cert-manager cert-manager.io/disable-validation=true  
  
kubectl apply --validate=false -f https://github.com/jetstack/cert-manager/releases/download/v1.0.1/cert-manager.yaml
```

9. Cert manager will need a custom ClusterIssuer resource to handle requesting SSL certificates.

```
code clusterissuer.yml
```

The following resource configuration should work as is:

```
apiVersion: cert-manager.io/v1  
kind: ClusterIssuer  
metadata:  
  name: letsencrypt-prod  
spec:  
  acme:  
    # The ACME server URL  
    server: https://acme-v02.api.letsencrypt.org/directory  
    # Email address used for ACME registration  
    email: user@fabmedical.com  
    # Name of a secret used to store the ACME account private key  
    privateKeySecretRef:  
      name: letsencrypt-prod  
    # Enable HTTP01 validations  
    solvers:  
    - http01:  
        ingress:  
          class: nginx
```

10. Save changes and close the editor.

11. Create the issuer using `kubectl`.

```
kubectl create --save-config=true -f clusterissuer.yml
```

12. Now you can create a certificate object.

Note:

Cert-manager might have already created a certificate object for you using ingress-shim.

To verify that the certificate was created successfully, use the `kubectl describe certificate tls-secret` command.

If a certificate is already available, skip to step 16.

```
code certificate.yml
```

Use the following as the contents and update the `[SUFFIX]` and `[AZURE-REGION]` to match your ingress DNS name:

```
apiVersion: cert-manager.io/v1
kind: Certificate
metadata:
  name: tls-secret
spec:
  secretName: tls-secret
  dnsNames:
    - fabmedical-[SUFFIX]-ingress.[AZURE-REGION].cloudapp.azure.com
  issuerRef:
    name: letsencrypt-prod
    kind: ClusterIssuer
```

13. Save changes and close the editor.

14. Create the certificate using `kubectl`.

```
kubectl create --save-config=true -f certificate.yml
```

Note: To check the status of the certificate issuance, use the `kubectl describe certificate tls-secret` command and look for an *Events* output similar to the following:

Type	Reason	Age	From	Message
Normal	Generated	38s	cert-manager	Generated new private key
Normal	GenerateSelfSigned	38s	cert-manager	Generated temporary self signed certificate
Normal	OrderCreated	38s	cert-manager	Created Order resource "tls-secret-3254248695"
Normal	OrderComplete	12s	cert-manager	Order "tls-secret-3254248695" completed successfully
Normal	CertIssued	12s	cert-manager	Certificate issued successfully

It can take between 5 and 30 minutes before the `tls-secret` becomes available. This is due to the delay involved with provisioning a TLS cert from letsencrypt.

15. Now you can create an ingress resource for the content applications.

```
code content.ingress.yml
```

Use the following as the contents and update the [SUFFIX] and [AZURE-REGION] to match your ingress DNS name:

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: content-ingress
  annotations:
    kubernetes.io/ingress.class: nginx
    certmanager.k8s.io/cluster-issuer: letsencrypt-prod
    nginx.ingress.kubernetes.io/rewrite-target: /$1
spec:
  tls:
    - hosts:
        - fabmedical-[SUFFIX]-ingress.[AZURE-REGION].cloudapp.azure.com
      secretName: tls-secret
  rules:
    - host: fabmedical-[SUFFIX]-ingress.[AZURE-REGION].cloudapp.azure.com
      http:
        paths:
          - path: /(.*)
            backend:
              serviceName: web
              servicePort: 80
          - path: /content-api/(.*)
            backend:
              serviceName: api
              servicePort: 3001
```

16. Save changes and close the editor.

17. Create the ingress using `kubectl`.

```
kubectl create --save-config=true -f content.ingress.yml
```

18. Refresh the ingress endpoint in your browser. You should be able to visit the speakers and sessions pages and see all the content.

19. Visit the api directly, by navigating to `/content-api/sessions` at the ingress endpoint.

```
[{"_id": "54b321f979cfa6002dd73477", "speakerNames": ["Theresa Zesiewicz", "Kevin Allison", "Israt Jahan", "Jessica Shaw", "F. Reed Murtagh", "Tracy Jones", "Clifton Gooch", "Jason Salemi", "Matthew B. Klein", "Guy Miller", "Kelly Sullivan"], "speakers": [], "trackNames": ["Visual Studio/Azure intersection", "ASP.NET / HTML5 intersection"], "tracks": ["124", "125"]}]
```

20. Test TLS termination by visiting both services again using `https`.

It can take between 5 and 30 minutes before the SSL site becomes available. This is due to the delay involved with provisioning a TLS cert from letsencrypt.

After the hands-on lab

Duration: 10 mins

In this exercise, you will de-provision any Azure resources created in support of this lab.

1. Delete the Resource Groups in which you placed all your Azure resources.
 - From the Portal, navigate to the blade of your **Resource Group** and then select **Delete** in the command bar at the top.
 - Confirm the deletion by re-typing the resource group name and selecting **Delete**.
2. Delete the Service Principal created on Task 3: Create a Service Principal before the hands-on lab.

```
az ad sp delete --id "Fabmedical-sp"
```

You should follow all steps provided *after* attending the Hands-on lab.